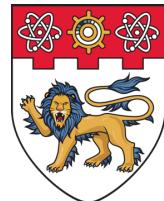


# *Convolutional Neural Networks on Graphs*

Xavier Bresson

School of Computer Science and Engineering  
Nanyang Technological University (NTU)



**NANYANG  
TECHNOLOGICAL  
UNIVERSITY  
SINGAPORE**

**NATIONAL RESEARCH FOUNDATION  
PRIME MINISTER'S OFFICE  
SINGAPORE**

Joint work with M. Defferrard (EPFL), P. Vandergheynst (EPFL),  
M. Bronstein (USI), F. Monti (USI), R. Levie (TAU)

Isaac Newton Institute  
Generative Models, Parameter Learning and Sparsity  
Nov 1<sup>st</sup> 2017

# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- Spectral Graph Theory
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- Spectral ConvNets
  - *SplineNets*
  - *ChebNets\** [NIPS'16]
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\** [NIPS'17]
- Conclusion

# Outline

## ➤ Deep Learning

- *A brief history*
- *High-dimensional learning*

## ➤ Convolutional Neural Networks

- *Architecture*
- *Non-Euclidean Data*

## ➤ Spectral Graph Theory

- *Graphs and operators*
- *Spectral decomposition*
- *Fourier analysis*
- *Convolution*
- *Coarsening*

## ➤ Spectral ConvNets

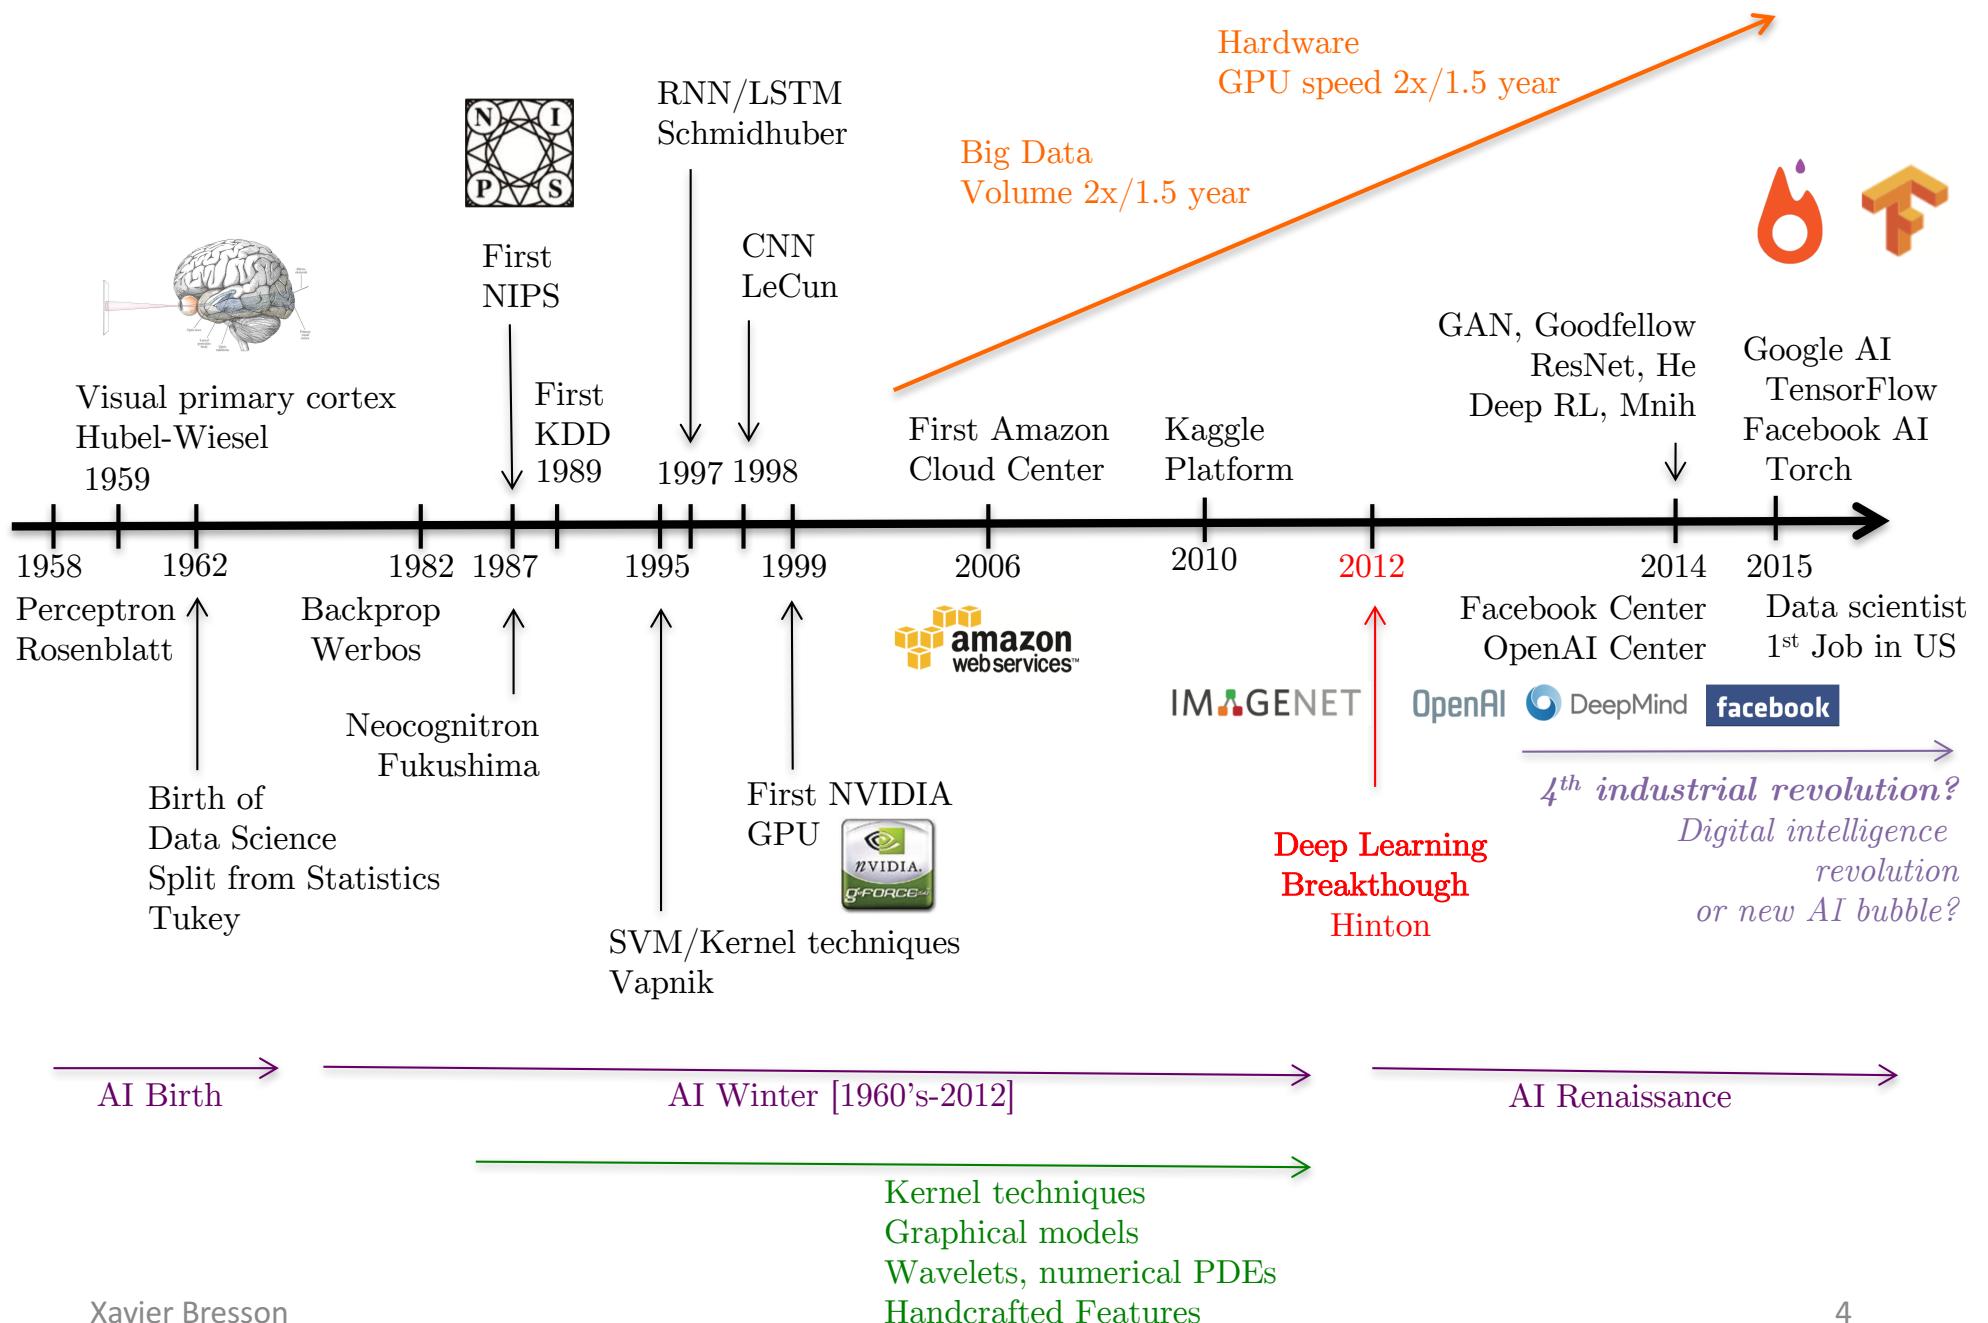
- *SplineNets*
- *ChebNets\* [NIPS'16]*
- *GraphConvNets*
- *CayleyNets\**

## ➤ Spectral ConvNets on Multiple Graphs

- *Recommender Systems\* [NIPS'17]*

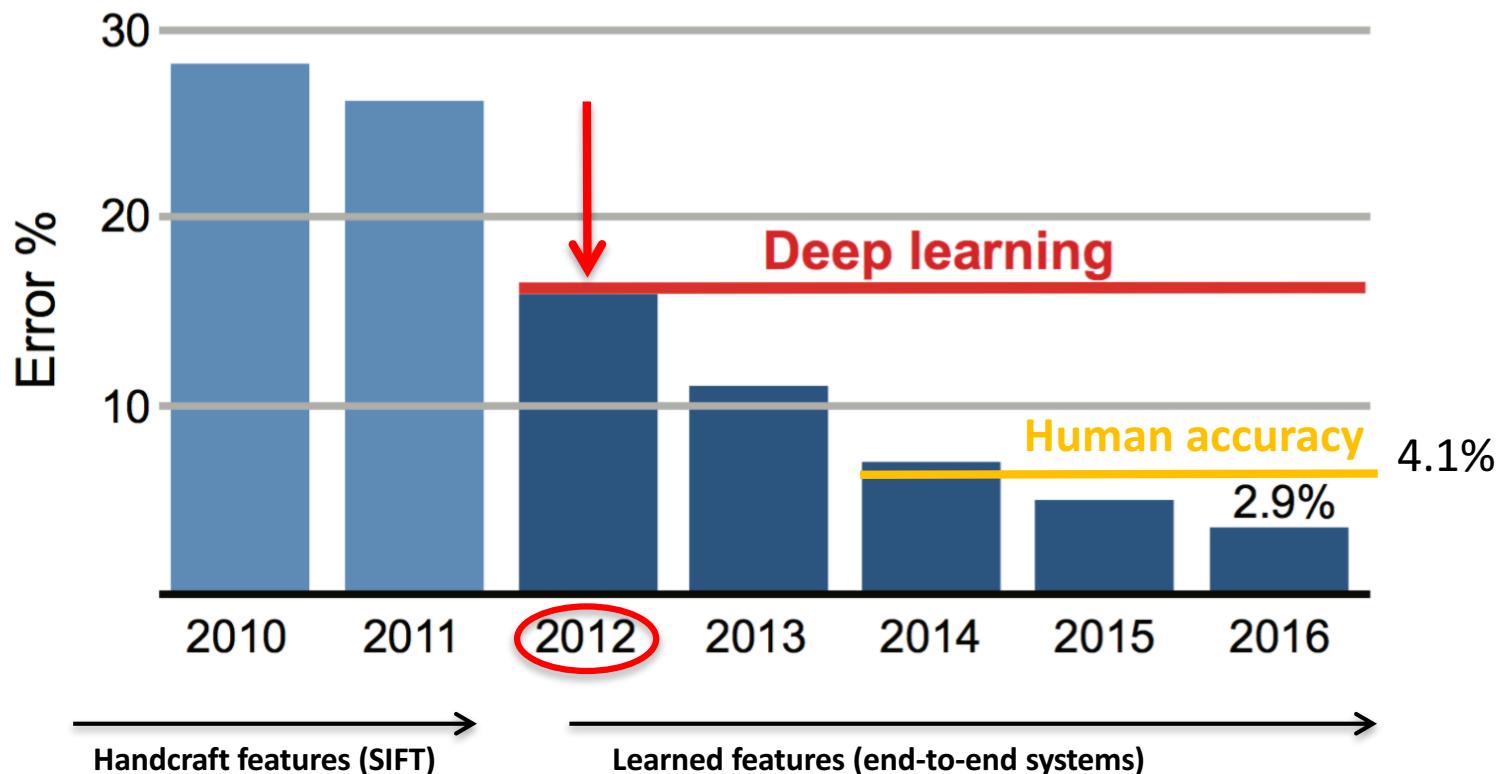
## ➤ Conclusion

# A brief history of artificial neural networks

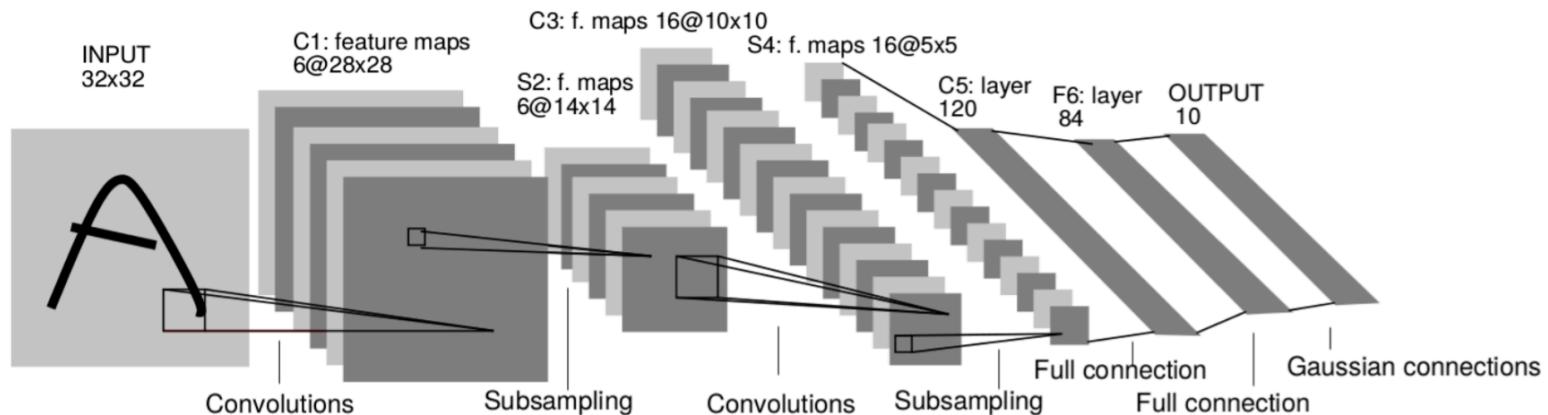


# Breakthrough in image recognition

IMAGENET



# Convolutional neural networks: LeNet-5<sup>[1]</sup>

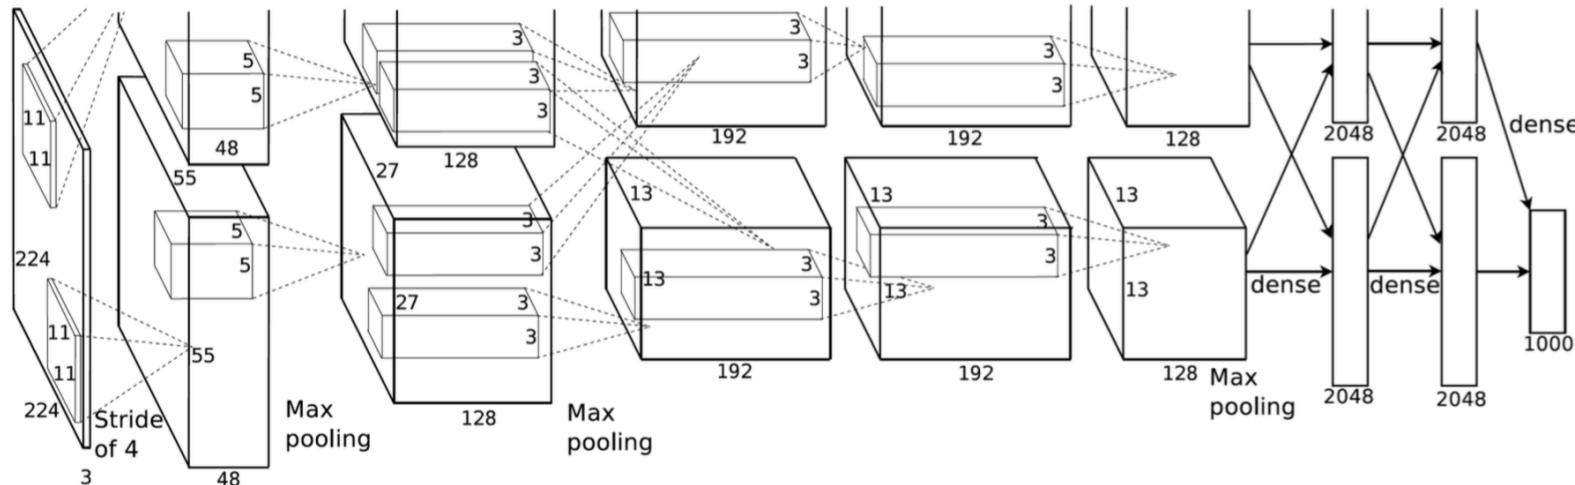


- 2 convolutional + 2 fully connected layers
- tanh non-linearity
- 1M parameters
- Training set: MNIST 70K images
- Trained on CPU

1998

[1] LeCun et al. 1998

# Convolutional neural networks: AlexNet<sup>[2]</sup>



- 5 convolutional + 3 fully connected layers
- ReLU non-linearity
- Dropout regularization
- 60M parameters
- Trained set: ImageNet 1.5M images
- Trained on GPU

2012

More learning capacity

More data

More computational power

[2] Krizhevsky, Sutskever, Hinton 2012

# Outline

## ➤ Deep Learning

- *A brief history*
- *High-dimensional learning*

## ➤ Convolutional Neural Networks

- *Architecture*
- *Non-Euclidean Data*

## ➤ Spectral Graph Theory

- *Graphs and operators*
- *Spectral decomposition*
- *Fourier analysis*
- *Convolution*
- *Coarsening*

## ➤ Spectral ConvNets

- *SplineNets*
- *ChebNets\* [NIPS'16]*
- *GraphConvNets*
- *CayleyNets\**

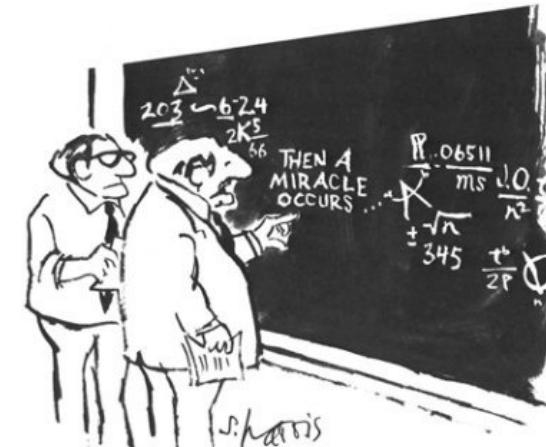
## ➤ Spectral ConvNets on Multiple Graphs

- *Recommender Systems\* [NIPS'17]*

## ➤ Conclusion

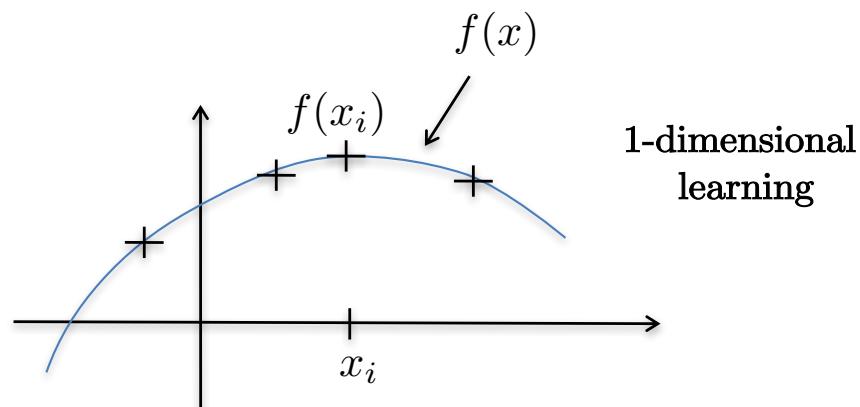
# High-dimensional learning

- NNs are powerful to solve high-dimensional learning problems.



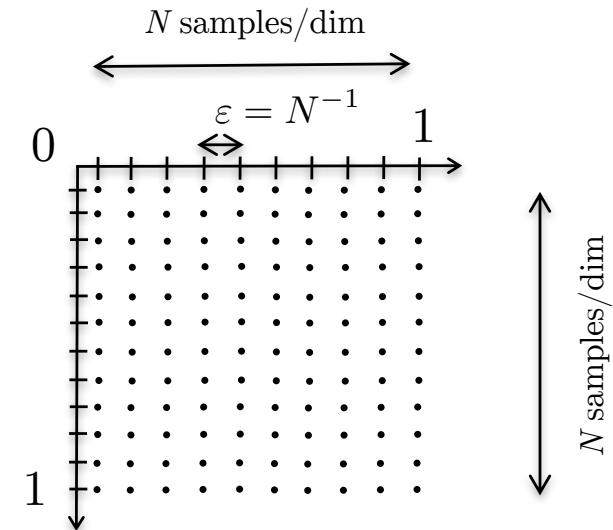
"I think you should be more explicit here in step two."

- Prototypal learning: Evaluate function  $f(x)$  given  $n$  samples  $\{x_i, f(x_i)\}$ . How? Suppose  $f$  is **regular**, use **interpolation** of close points.



# Curse of dimensionality

- What about in high dimensions? No interpolation possible because we can never have enough close points to make a good interpolation.
- How many points to cover the 2D plane?  $\epsilon^{-2}$   
 $\Rightarrow d$ -dim cube  $[0,1]^d$  requires  $\epsilon^{-d}$  points.
- $\text{dim(image)} = 512 \times 512 \approx 10^6$   
For  $N=10$  samples/dim  $\Rightarrow 10^{1,000,000}$  points!
- Curse of dim:
  - :( The number of  $x$  of  $f$  is exponential.
  - :( All samples are far from each other in high-dim<sup>[3]</sup>.
  - :( Optimization: Also need  $\epsilon^{-d}$  points to find solution.

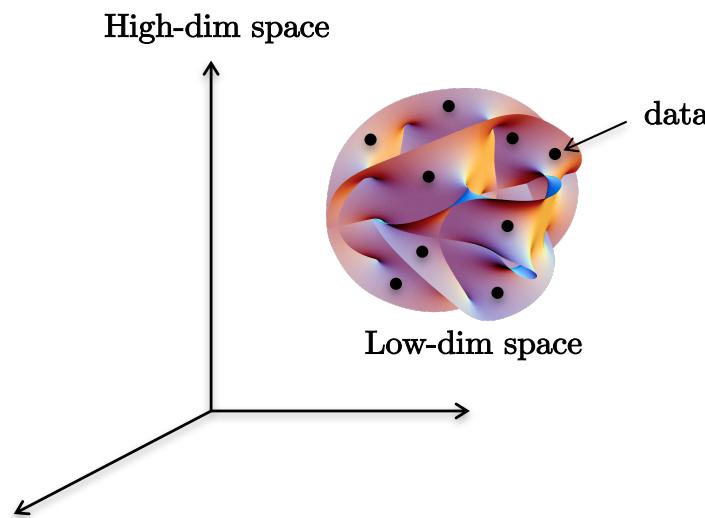


$$\min_f L(f) = \mathbb{E}[\ell(f(x_i), y_i)]$$

[3] Beyer 1998

# Blessing of structure

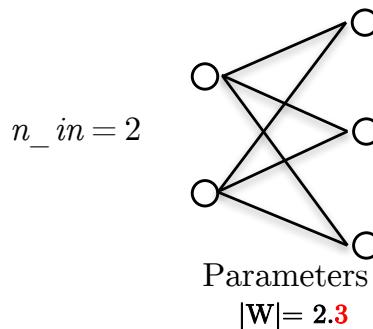
- In our universe:
  - ☺ Data have **structures**, they concentrate in (non-)smooth **low-dim spaces**.



- ☺ These structures can be expressed as **mathematical properties** (statistics, geometry, groups, etc).

# Learning exponential functions

- Neural networks can learn **parametric** exponential functions with **linear** number of parameters.



$$\begin{aligned}n_{in} &= 2 \\n_{out} &= 3 \\ \# \text{configurations} &= \\ 2^{n_{out}} &= 2^3 = 8\end{aligned}$$

⇒ Number of parameters grow linearly,  
but function capacity grows exponentially.

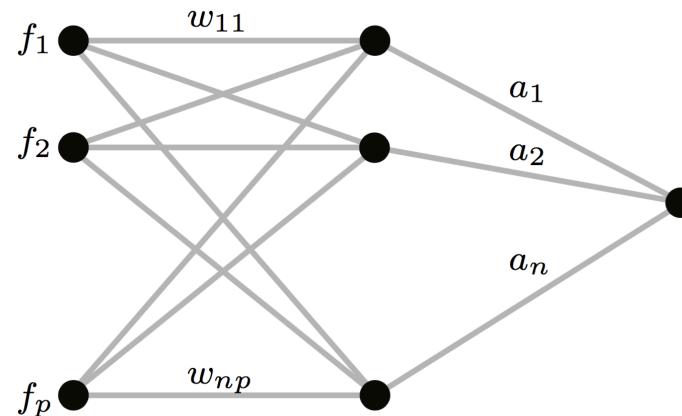
- A very large number of parameters can be learned by backpropagation, order of billions.  
⇒ Huge capacity to learn complex data<sup>[4]</sup>.

[4] Zhang, Bengio, Hardt, Recht, Vinyals, 2016

# Universal function approximation<sup>[4]</sup>

**Universal Approximation Theorem** Let  $\xi$  be a non-constant, bounded, and monotonically-increasing continuous activation function,  $y : [0, 1]^p \rightarrow \mathbb{R}$  continuous function, and  $\epsilon > 0$ . Then,  $\exists n$  and parameters  $\mathbf{a}, \mathbf{b} \in \mathbb{R}^n$ ,  $\mathbf{W} \in \mathbb{R}^{n \times p}$  s.t.

$$\left| \sum_{i=1}^n a_i \xi(\mathbf{w}_i^\top \mathbf{f} + b_i) - y(\mathbf{f}) \right| < \epsilon \quad \forall \mathbf{f} \in [0, 1]^p$$



[5] Cybenko 1989, Hornik 1991

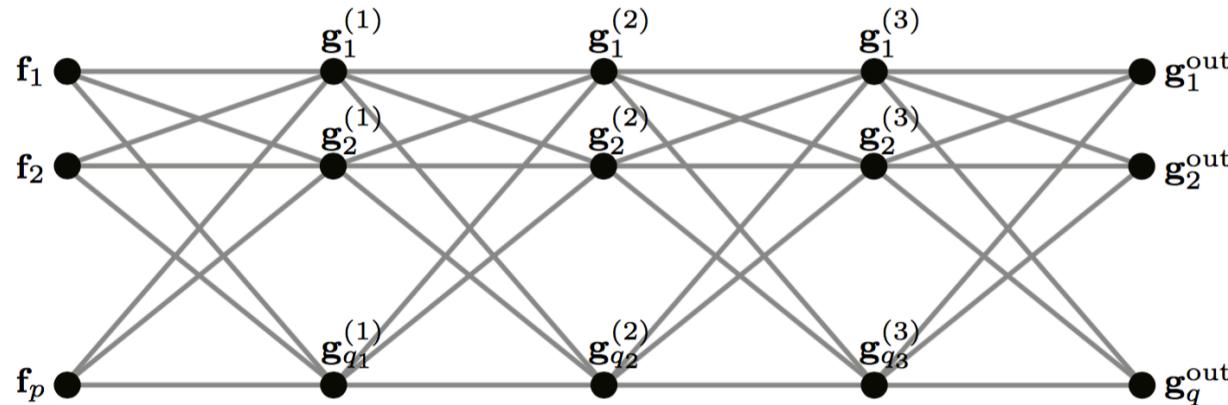
# Shallow neural nets as universal approximators?

- ☺ Any continuous function can be approximated arbitrarily well by a neural network with a single hidden layer. Then,
  - How many neurons?
  - How long is the optimization?
  - Does it generalize well/overfit?
- ☹ Shallow NNs: Data (s.a. image) requires NNs with an exponential number of neurons for a single layer, optimizes badly, and does not generalize.
- ☺ Deep NNs: Multiple layers and data structure can alleviate the above issues<sup>[6]</sup>.

[6] Montufar, Pascanu, Cho, Bengio 2014

# Fully connected networks

$$\begin{aligned}\mathbf{f}_l &= l\text{-th image feature (R,G,B channels), } \dim(\mathbf{f}_l) = n \times 1 \\ \mathbf{g}_l^{(k)} &= l\text{-th feature map, } \dim(\mathbf{g}_l^{(k)}) = n_l^{(k)} \times 1\end{aligned}$$



Deep neural network consists of multiple layers.

Linear layer

$$\mathbf{g}_l^{(k)} = \xi \left( \sum_{l'=1}^{q_{k-1}} \mathbf{W}_{l,l'}^{(k)} \xi \left( \sum_{l'=1}^{q_{k-2}} \mathbf{W}_{l,l'}^{(k-1)} \xi \left( \dots \mathbf{f}_{l'} \right) \right) \right)$$

Activation, e.g.  $\xi(x) = \max\{x, 0\}$  rectified linear unit (ReLU)

# Fully connected networks

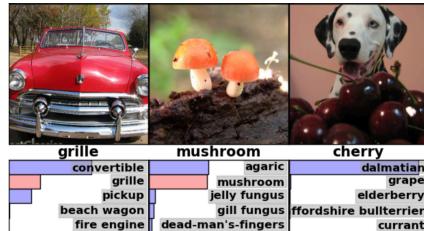
- FC networks capture non-linear **factorizations** of data.
- They **overfit**: Too many parameters to learn.
- NN architectures require additional data **structures** to successfully analyze image, sound, document, etc.

# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- Spectral Graph Theory
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- Spectral ConvNets
  - *SplineNets*
  - *ChebNets\* [NIPS'16]*
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\* [NIPS'17]*
- Conclusion

# Convolutional neural networks<sup>[1]</sup>

- Data (image, video, sound) are **compositional**, i.e. they are formed of patterns that are:
  - Local
  - Stationary
  - Multi-scale (/hierarchical)
- ConvNets leverage the compositionality structure: They extract compositional features and feed them to classifier, recommender, etc (end-to-end).

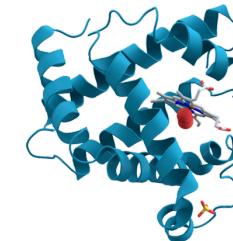


Computer Vision

```
/* Duplicate LSM field information. The lsm_rule is opaque, so
 * we can't copy it directly.
 * static inline int audit_dupe_lsm_field(struct audit_field *df,
 *                                     struct audit_field *sf)
```

```
    int ret = 0;
    char lsm_str[1];
    /* Don't need to copy df.lsm_str */
    lsm_str = kstrdup(sf->lsm_str, GFP_KERNEL);
    if (!lsm_str)
        return -ENOMEM;
    sf->lsm_str = lsm_str;
    /* Return a copy of lsm_rule */
    ret = security_audit_rule_init(df->xtype, df->op, df->lsm_str,
                                   (void *) &df->lsm_rule);
    /* Keep current rule valid until a new one is found in case they
     * become valid after a policy reload. */
    if (ret == -EINVAL) {
        /* If the old rule for LSM '\n' is invalid... */
        df->lsm_str = NULL;
    }
    return ret;
```

NLP



Drug discovery

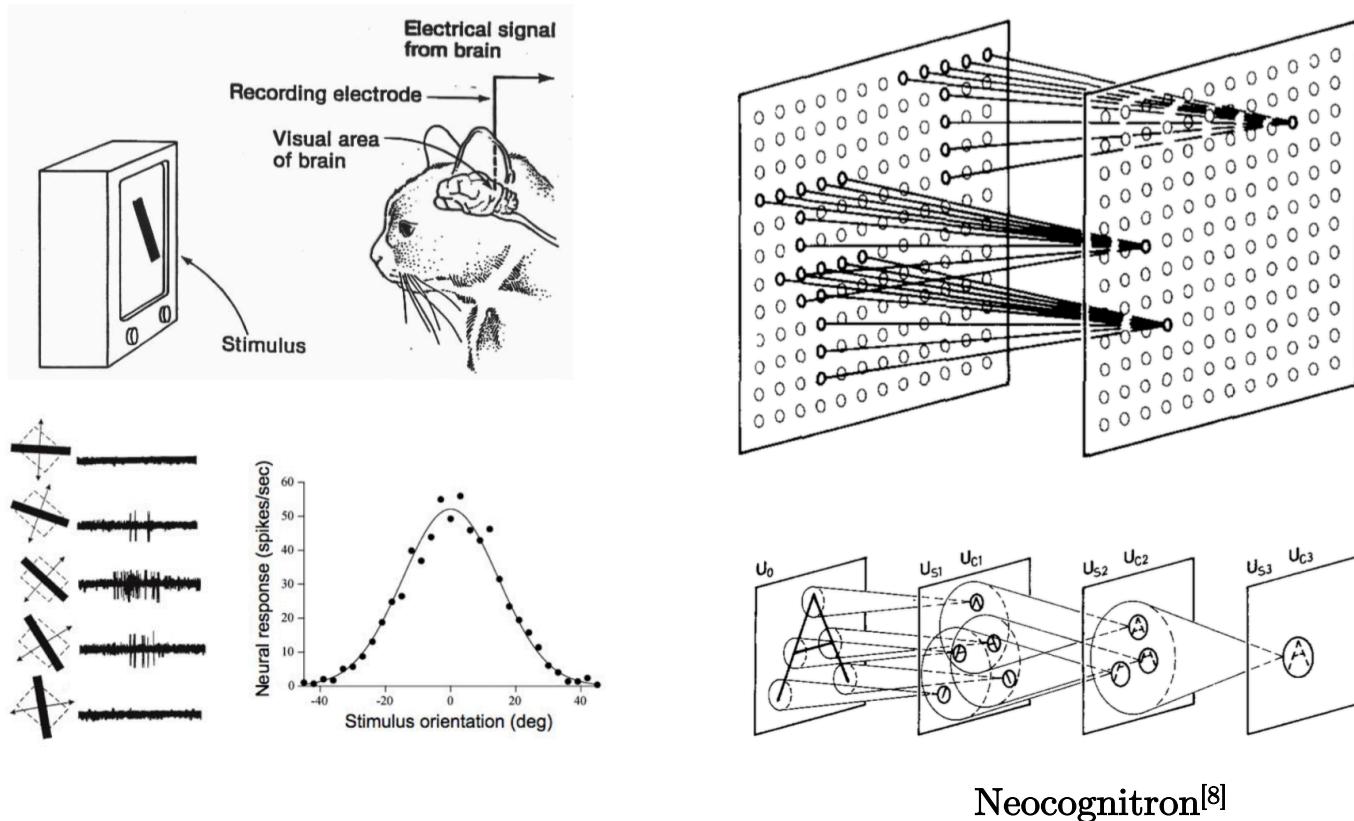


Games

[1] LeCun et al. 1998

# Key properties

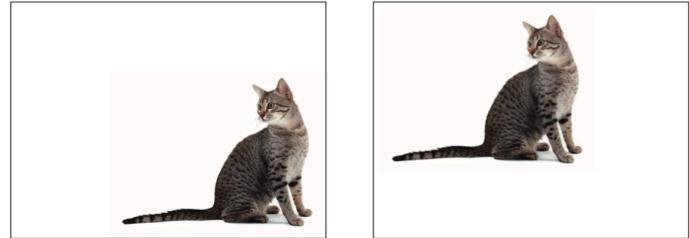
- **Locality:** Property inspired by visual cortex neurons<sup>[7]</sup>.
- **Local receptive fields** activate in the presence of (learned) local features.



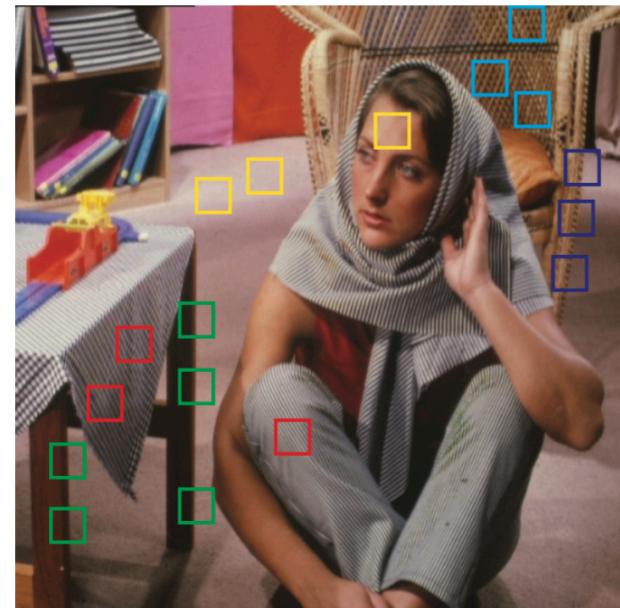
[7] Cybenko 1989, Hornik 1991, [8] Fukushima 1980

# Key properties

- **Stationarity**  $\Leftrightarrow$  Translation invariance

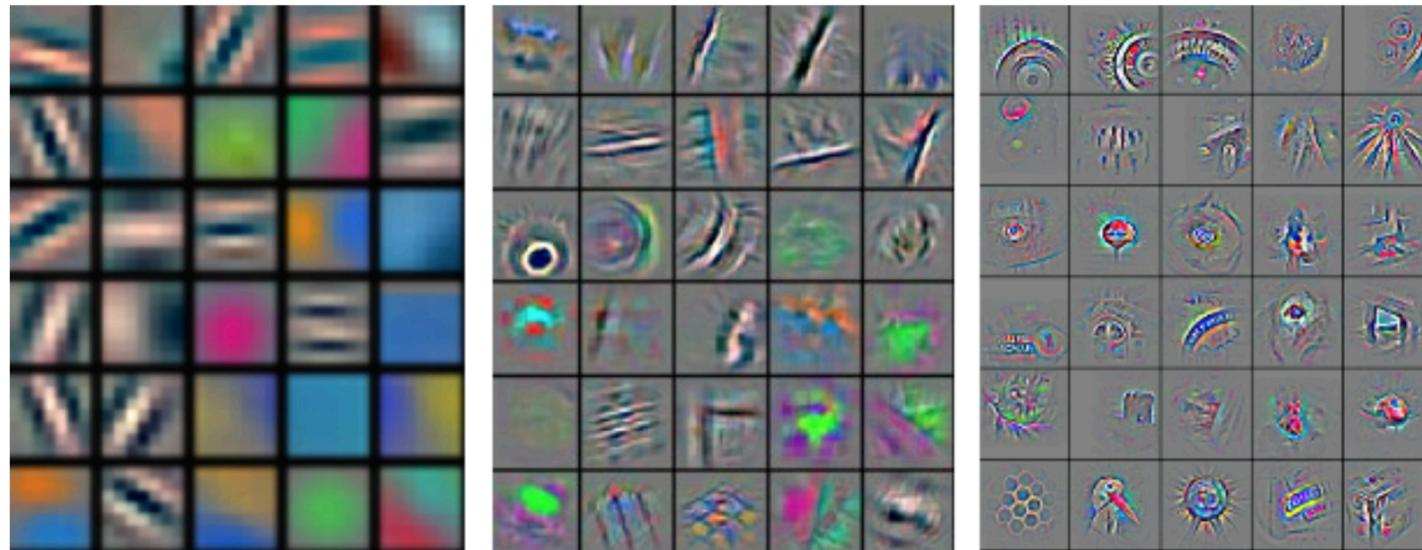


- **Local stationarity**  $\Leftrightarrow$  Similar patches shared across the data domain.



# Key properties

- **Multi-scale:** Simple structures combine to compose slightly more abstract structures, and so on, in a hierarchical way.
- Also inspired by brain visual primary cortex (V1 and V2 neurons).

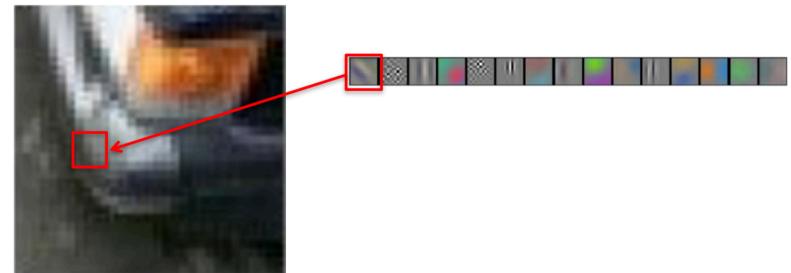


Features learned by a ConvNet become increasingly complex at deeper layers.<sup>[9]</sup>

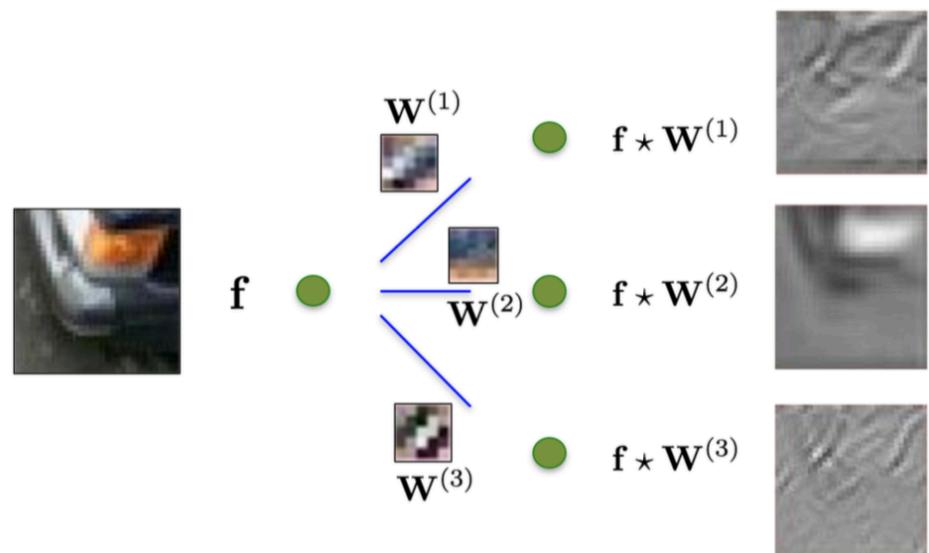
[9] Zeiler, Fergus 2013

# Implementation

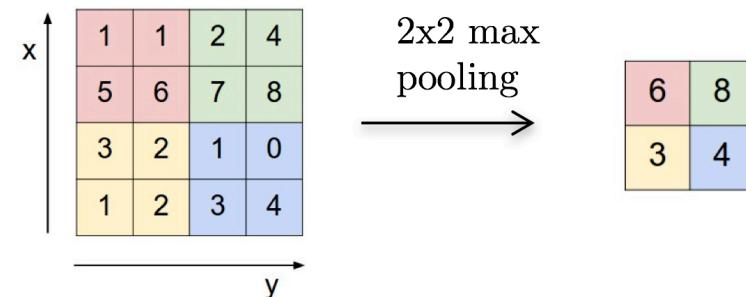
- **Locality:** compact support kernels  
 $\Rightarrow O(1)$  parameters per filter.



- **Stationarity:** convolutional operators  
 $\Rightarrow O(n \log n)$  in general (FFT) and  
 $O(n)$  for compact kernels.

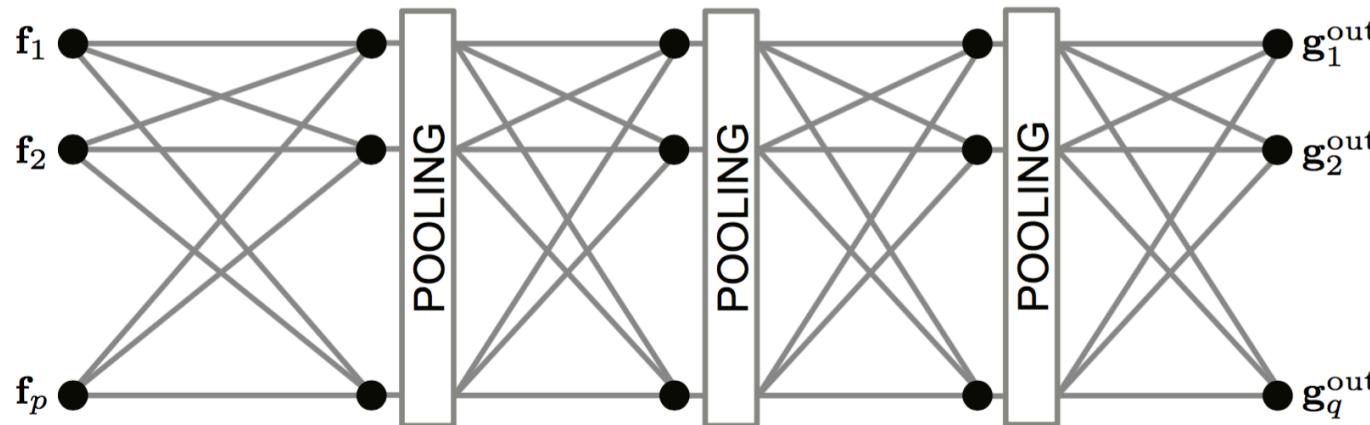


- **Multi-scale:** downsampling + pooling  $\Rightarrow O(n)$



# Compositional features

$$\begin{aligned}\mathbf{f}_l &= l\text{-th image feature (R,G,B channels), } \dim(\mathbf{f}_l) = n \times 1 \\ \mathbf{g}_l^{(k)} &= l\text{-th feature map, } \dim(\mathbf{g}_l^{(k)}) = n_l^{(k)} \times 1\end{aligned}$$



Compositional features consist of multiple convolutional + pooling layers.

**Linear layer**

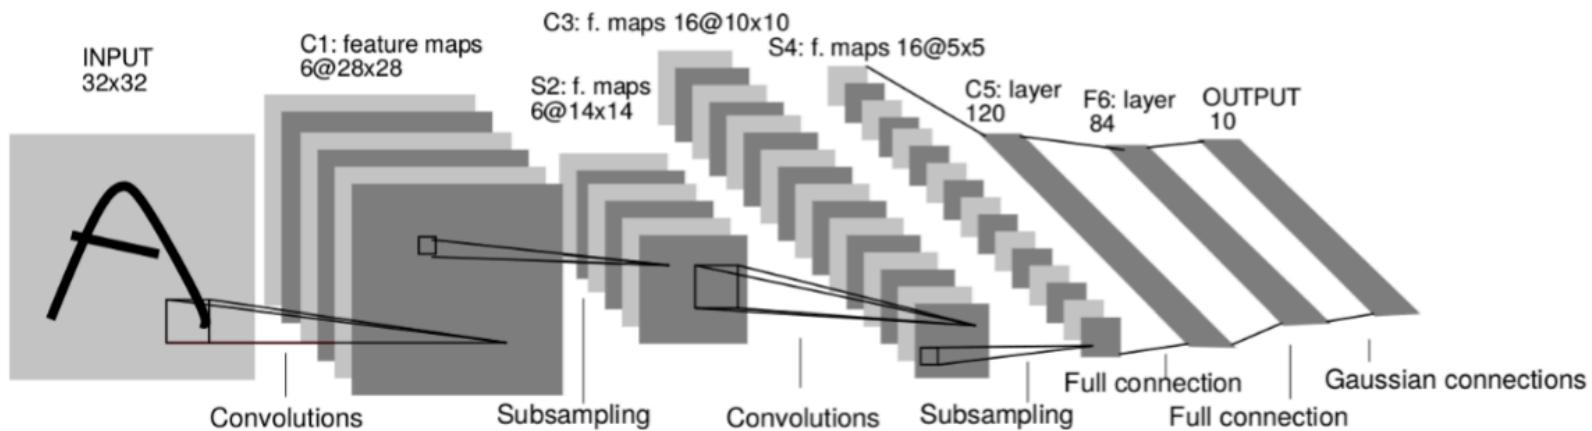
$$\mathbf{g}_l^{(k)} = \xi \left( \sum_{l'=1}^{q_k-1} \mathbf{W}_{l,l'}^{(k)} \star \xi \left( \sum_{l'=1}^{q_{k-1}} \mathbf{W}_{l,l'}^{(k-1)} \star \xi \left( \dots \mathbf{f}_{l'} \right) \right) \right)$$

**Activation**, e.g.  $\xi(x) = \max\{x, 0\}$  rectified linear unit (ReLU)

**Pooling**

$$\mathbf{g}_l^{(k)}(x) = \|\mathbf{g}_l^{(k-1)}(x') : x' \in \mathcal{N}(x)\|_p \quad p = 1, 2, \text{ or } \infty$$

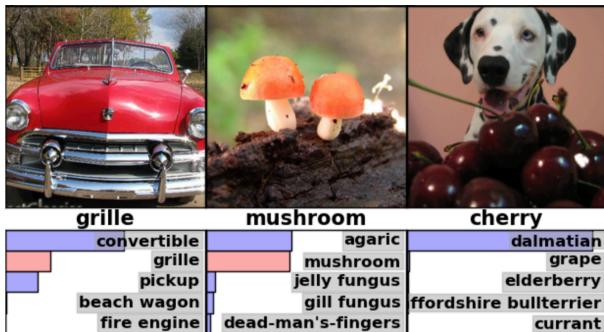
# ConvNets<sup>[1]</sup>



- ☺ Filters localized in space (Locality)
- ☺ Convolutional filters (Stationarity)
- ☺ Multiple layers (Multi-scale)
- ☺  $O(1)$  parameters per filter (independent of input image size  $n$ )
- ☺  $O(n)$  complexity per layer (filtering done in the spatial domain)

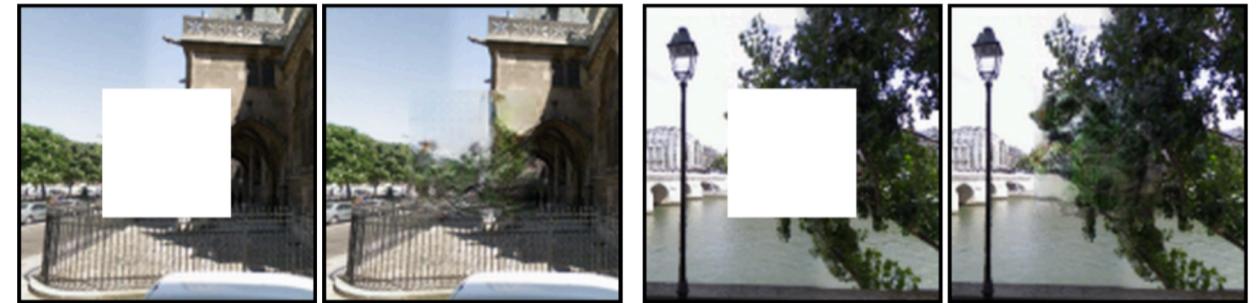
[1] LeCun et al. 1998

# ConvNets in computer vision



*Image/object recognition*

[Krizhevsky-Sutskever-Hinton'12]



[Pathak-Efros-etal'16]



*Image captioning*

[Karpathy-Li'15]



⇒ Highly successful - ConvNets is the Swiss knife of Computer Vision!

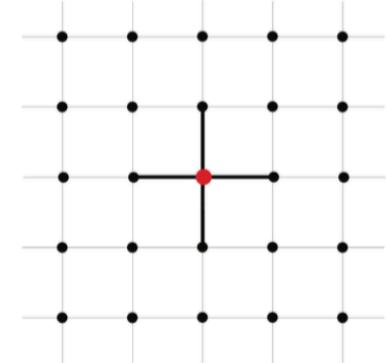
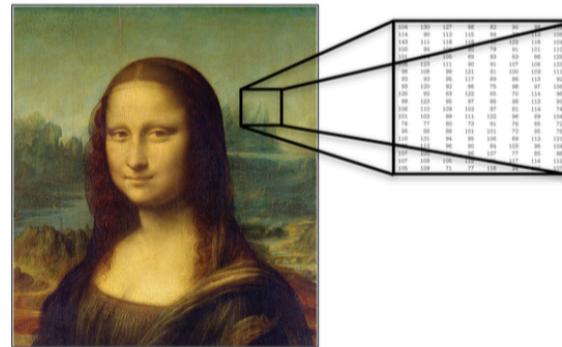


# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- Spectral Graph Theory
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- Spectral ConvNets
  - *SplineNets*
  - *ChebNets\* [NIPS'16]*
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\* [NIPS'17]*
- Conclusion

# Data domain for ConvNets

- Image, volume, video: 2D, 3D, 2D+1 Euclidean domains



2D grids

- Sentence, word, sound: 1D Euclidean domain

O ROMEO, ROMEO, WHEREFORE ART THOU ROMEO?  
Although we use "wherefore," if at all, as a synonym for "why," Juliet uses the word in a more limited sense. By "wherefore?" Juliet means "for what purpose?" If she had merely asked "Why art thou Romeo?" she would be distinguishing between two types of causes: "why" —from what cause? (in the past) and "for what purpose" (in the future). "Wherefore" clearly emphasizes the latter sense, which is why "why" and "wherefores" are different things.

"Wherefore" and its partner "therefore" reflect the basic tendency of English to use spatial ideas—“where,” “there”—to represent logical ideas, such as cause and effect.

WHAT'S IN A NAME? THAT WHICH WE CALL A ROSE BY ANY OTHER WORD WOULD SMELL AS SWEET  
If there's such a thing as generic Shakespeare today, this is it. Both "sweet" and "rose" are instant Bard, although the latter is, as many forget, merely a paraphrase. From the romantic declamation to the crass advertisement, these phrases have served generations with complete flexibility.

"What's in a name?" is the less specific of the two phrases, and also the less common. Juliet here merely rehearses in a different form the point of "what's in a name": more or less like a game of telephone, from the partner to the generic. Now, on general, she insists, ought to be separable from the things they name. Romeo never does change his name, and it wouldn't have done much good anyway. Whether or not he's essentially a Montague, and Juliet essentially a Capulet, their families will continue to act that way.

"That which we call a rose/ By any other word would smell as sweet" seems blindingly obvious now, but we've accustomed to the paraphrase, which never occurs to the playwright or his audience. It's a little futile to second-guess Shakespeare now, but he did have to fill out a line and a half of blank verse. Regarding Juliet's use of "word" instead of "name," we can perhaps be grateful; she already uses "name" six times in fifteen and a half lines.

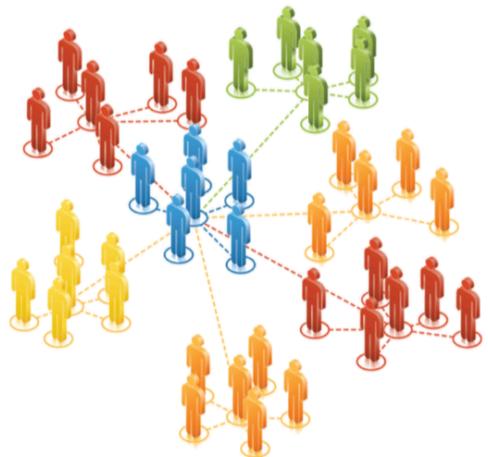
978 114 90 112 115 94 76 113 101



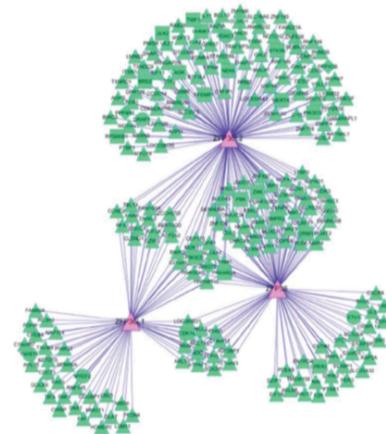
1D grid

- These domains have nice regular spatial structures.  
⇒ All CNN operations are math well defined and fast (convolution, pooling).

# Non-Euclidean data

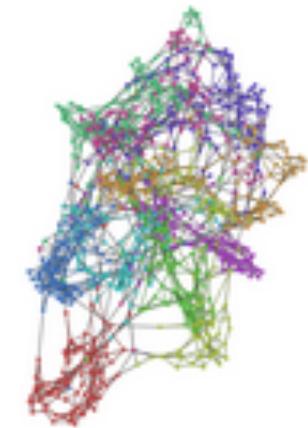


Social networks

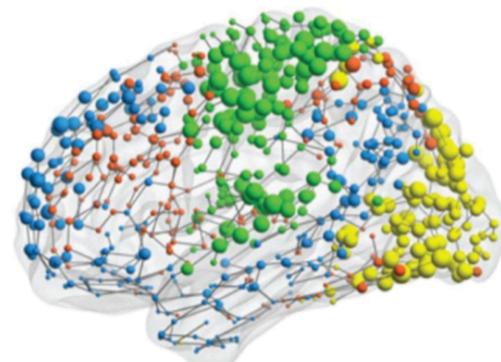


Regulatory networks

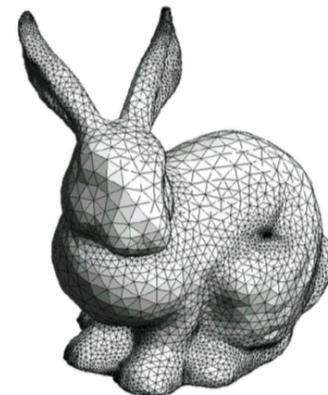
=



Graphs/  
Networks



Functional networks



3D shapes

- Also NLP, physics, social science, communication networks, etc.

# Challenges of graph deep learning

- Extend neural network techniques to **graph-structured data**.
- **Assumption:** Non-Euclidean data are locally stationary and manifest hierarchical structures.
- How to define **compositionality on graphs**? (convolution and pooling on graphs)
- How to make them **fast**? (linear complexity)

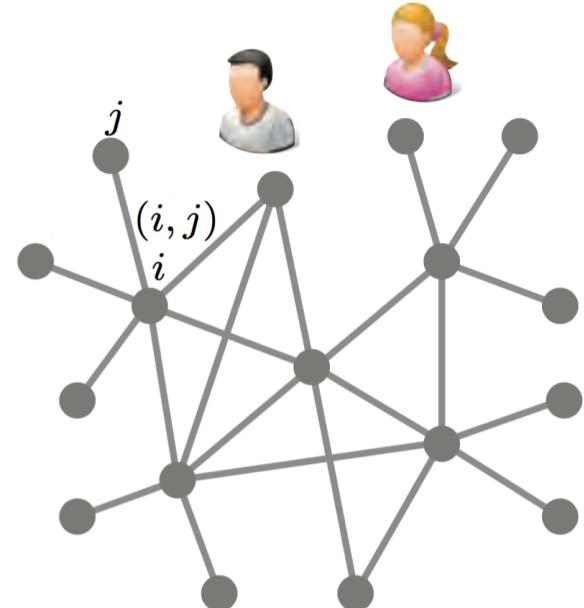
# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- **Spectral Graph Theory**
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- Spectral ConvNets
  - *SplineNets*
  - *ChebNets*
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\* [NIPS'17]*
- Conclusion

# Graphs<sup>[10]</sup>

- **Graph**  $\mathcal{G} = (\mathcal{V}, \mathcal{E})$
- **Vertices**  $\mathcal{V} = \{1, \dots, n\}$
- **Edges**  $\mathcal{E} \subseteq \mathcal{V} \times \mathcal{V}$
- **Vertex weights**  $b_i > 0$  for  $i \in \mathcal{V}$
- **Edge weights**  $a_{ij} \geq 0$  for  $(i, j) \in \mathcal{E}$
- **Vertex fields**  $L^2(\mathcal{V}) = \{f : \mathcal{V} \rightarrow \mathbb{R}^h\}$   
**Represented as**  $\mathbf{f} = (f_1, \dots, f_n)$
- **Hilbert space with inner product**

$$\langle f, g \rangle_{L^2(\mathcal{V})} = \sum_{i \in \mathcal{V}} a_i f_i g_i$$



[10] Chung 1994

# Calculus on graphs

- **Gradient operator**  $\nabla : L^2(\mathcal{V}) \rightarrow L^2(\mathcal{E})$

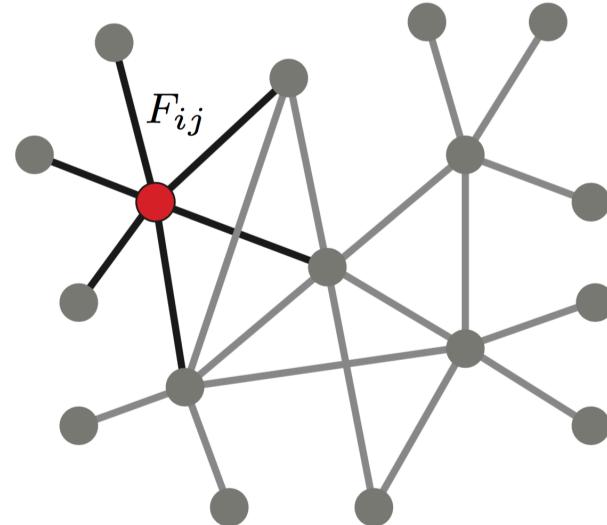
$$(\nabla f)_{ij} = \sqrt{a_{ij}}(f_i - f_j)$$

- **Divergence operator**  $\operatorname{div} : L^2(\mathcal{E}) \rightarrow L^2(\mathcal{V})$

$$(\operatorname{div} F)_i = \frac{1}{b_i} \sum_{j:(i,j) \in \mathcal{E}} \sqrt{a_{ij}}(F_{ij} - F_{ji})$$

adjoint to the gradient operator

$$\langle F, \nabla f \rangle_{L^2(\mathcal{E})} = \langle \nabla^* F, f \rangle_{L^2(\mathcal{V})} = \langle -\operatorname{div} F, f \rangle_{L^2(\mathcal{V})}$$



# Graph Laplacian

- **Laplacian operator**  $\Delta : L^2(\mathcal{V}) \rightarrow L^2(\mathcal{V})$

$$(\Delta f)_i = \frac{1}{b_i} \sum_{j:(i,j) \in \mathcal{E}} a_{ij} (f_i - f_j)$$

difference between  $f$  and its local average (2<sup>nd</sup> derivative on graphs)

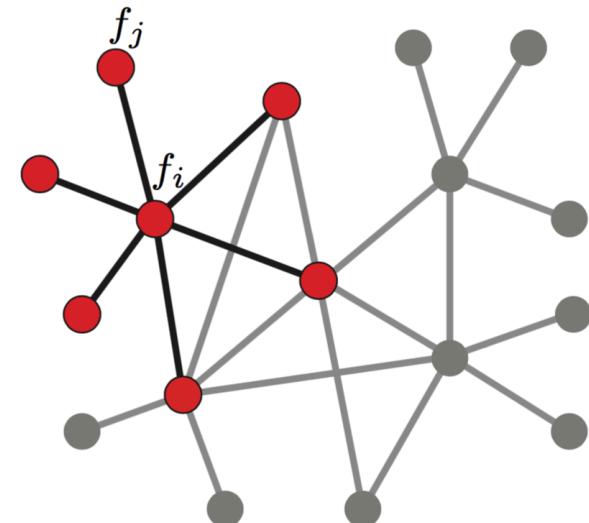
- **Core operator** in spectral graph theory.
- Represented as a **positive semi-definite**  $n \times n$  matrix

- **Unnormalized Laplacian**  $\Delta = \mathbf{D} - \mathbf{A}$

- **Normalized Laplacian**  $\Delta = \mathbf{I} - \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2}$

- **Random walk Laplacian**  $\Delta = \mathbf{I} - \mathbf{D}^{-1} \mathbf{A}$

where  $\mathbf{A} = (a_{ij})$  and  $\mathbf{D} = \text{diag}(\sum_{j \neq i} a_{ij})$



# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- **Spectral Graph Theory**
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- Spectral ConvNets
  - *SplineNets*
  - *ChebNets\* [NIPS'16]*
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\* [NIPS'17]*
- Conclusion

# Spectral decomposition

- A **Laplacian** of a graph of  $n$  vertices admits  **$n$  eigenvectors**

$$\Delta \phi_k = \lambda_k \phi_k, \quad k = 1, 2, \dots$$

- Eigenvectors are **real** and **orthonormal**  $\langle \phi_k, \phi_{k'} \rangle_{L^2(\mathcal{V})} = \delta_{kk'}$   
(due to self-adjointness)
- Eigenvalues are **non-negative**  $0 = \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$   
(due to positive-semidefiniteness)
- Eigendecomposition of a graph Laplacian

$$\Delta = \Phi^T \Lambda \Phi$$

where  $\Phi = (\phi_1, \dots, \phi_n)$  and  $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$

# Interpretation

- Find the **smoothest** orthonormal basis  $\Phi = (\phi_1, \dots, \phi_n)$  on a graph

$$\begin{aligned} \min_{\phi_k} E_{\text{Dir}}(\phi_k) \quad & \text{s.t.} \quad \|\phi_k\| = 1, \quad k = 2, 3, \dots, n \\ & \phi_k \perp \text{span}\{\phi_1, \dots, \phi_{k-1}\} \end{aligned}$$

where  $E_{\text{Dir}}$  is the **Dirichlet** energy = measure of smoothness of a function

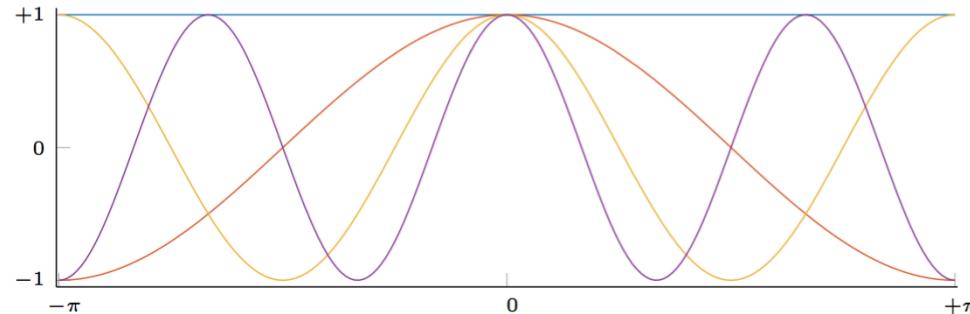
$$E_{\text{Dir}}(\mathbf{f}) = \mathbf{f}^\top \Delta \mathbf{f}$$

- **Solution:** first  $n$  Laplacian eigenvectors

$$\begin{aligned} \min_{\Phi \in \mathbb{R}^{n \times n}} \underbrace{\text{trace}(\Phi^\top \Delta \Phi)}_{\|\Phi\|_{\mathcal{G}} \text{ Dirichlet norm}} \quad & \text{s.t.} \quad \Phi^\top \Phi = \mathbf{I} \end{aligned}$$

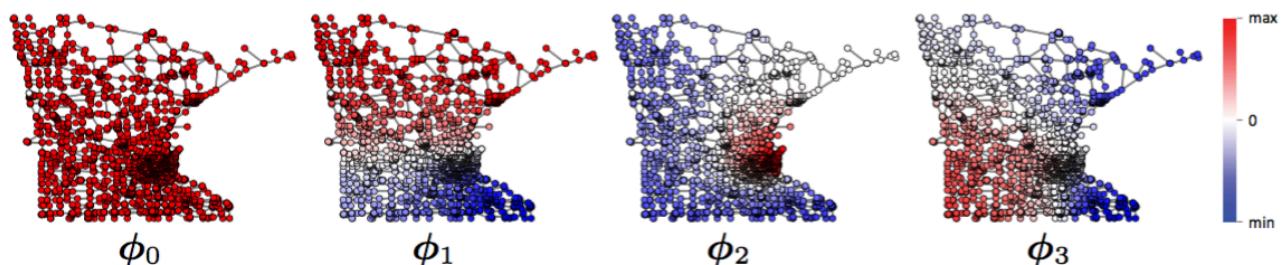
# Laplacian eigenvectors

- Euclidean domain:



First eigenvectors of 1D Euclidean Laplacian = standard Fourier basis

- Graph domain:



First Laplacian eigenvectors of a graph

Lap eigenvectors related to graph geometry  
(s.a. communities, hubs, etc), spectral clustering<sup>[10]</sup>

[10] Von Luxburg 2007

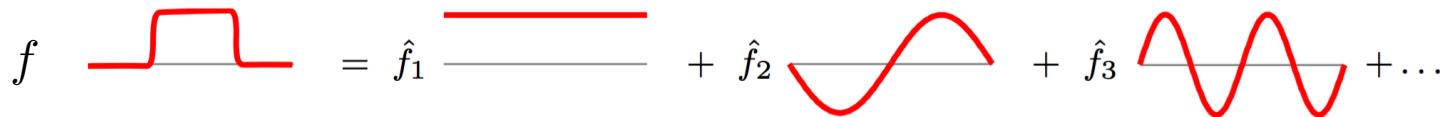
# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- **Spectral Graph Theory**
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- Spectral ConvNets
  - *SplineNets*
  - *ChebNets\* [NIPS'16]*
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\* [NIPS'17]*
- Conclusion

# Fourier analysis on Euclidean spaces

- A function  $f : [-\pi, \pi] \rightarrow \mathbb{R}$  can be written as Fourier series

$$f(x) = \sum_{k \geq 0} \underbrace{\frac{1}{2\pi} \int_{-\pi}^{\pi} f(x') e^{-ikx'} dx'}_{\hat{f}_k = \langle f, e^{-ikx} \rangle_{L^2([-\pi, \pi])}} e^{-ikx}$$



- Fourier basis  $e^{-ikx}$  = Laplace-Beltrami eigenfunctions:

$$-\Delta \phi_k = k^2 \phi_k$$

$$\begin{aligned}\phi_k &= \text{Fourier mode} \\ k &= \text{frequency of Fourier mode}\end{aligned}$$

# Fourier analysis on graphs<sup>[11]</sup>

- A function  $f : \mathcal{V} \rightarrow \mathbb{R}$  can be written as Fourier series

$$f_i = \sum_{k=1}^n \underbrace{\langle f, \phi_k \rangle_{L^2(\mathcal{V})}}_{\hat{f}_k} \phi_{k,i}$$

- $\hat{f}_k$  is the  $k$ -th graph Fourier coefficient.
- In matrix-vector notation, with the  $n \times n$  Fourier matrix  $\Phi = [\phi_1, \dots, \phi_n]$

$$\hat{\mathbf{f}} = \Phi^\top \mathbf{f} \quad \text{and} \quad \mathbf{f} = \Phi \hat{\mathbf{f}}$$

Fourier transform                          Inverse Fourier transform

- Graph Fourier basis  $\phi_k$  = Laplacian eigenvectors :

$$\Delta \phi_k = \lambda_k \phi_k$$

$$\begin{aligned}\phi_k &= \text{graph Fourier mode} \\ \lambda_k &= \text{(square) frequency}\end{aligned}$$

[11] Hammond, Vandergheynst, Gribonval, 2011

# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- **Spectral Graph Theory**
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- Spectral ConvNets
  - *SplineNets*
  - *ChebNets\* [NIPS'16]*
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\* [NIPS'17]*
- Conclusion

# Convolution on Euclidean spaces

- Given two functions  $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$  their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

- Shift-invariance:**  $f(x - x_0) \star g(x) = (f \star g)(x - x_0)$
- Convolution operator **commutes with Laplacian:**  $(\Delta f) \star g = \Delta(f \star g)$
- Convolution theorem:** Convolution can be computed in the Fourier domain as

$$\widehat{(f \star g)} = \hat{f} \cdot \hat{g}$$

- Efficient computation** using FFT:  $O(n \log n)$

# Convolution on discrete Euclidean spaces

- Convolution of two vectors  $\mathbf{f} = (f_1, \dots, f_n)^\top$  and  $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\begin{aligned}
 (\mathbf{f} \star \mathbf{g})_i &= \sum_m g_{(i-m) \text{ mod } n} \cdot f_m \\
 \mathbf{f} \star \mathbf{g} &= \underbrace{\left[ \begin{array}{ccccc} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{array} \right]}_{\text{Circulant matrix}} \left[ \begin{array}{c} f_1 \\ \vdots \\ f_n \end{array} \right]
 \end{aligned}$$

diagonalised by Fourier basis (Toeplitz)

$$= \Phi \begin{bmatrix} \hat{g}_1 & & & \\ & \ddots & & \\ & & & \hat{g}_n \end{bmatrix} \Phi^\top \mathbf{f} = \Phi(\Phi^\top \mathbf{g} \circ \Phi^\top \mathbf{f})$$

# Convolution on graphs<sup>[11]</sup>

- Spectral convolution of  $f, g \in L^2(\mathcal{V})$  can be defined by analogy

$$(f \star g)_i = \sum_{k \geq 1} \underbrace{\langle f, \phi_k \rangle_{L^2(\mathcal{V})} \langle g, \phi_k \rangle_{L^2(\mathcal{V})}}_{\text{product in the Fourier domain}} \phi_{k,i}$$

inverse Fourier transform

- In matrix-vector notation

$$\begin{aligned} \mathbf{f} \star \mathbf{g} &= \Phi (\Phi^\top \mathbf{g} \circ \Phi^\top \mathbf{f}) \\ &= \underbrace{\Phi \text{diag}(\hat{g}_1, \dots, \hat{g}_n) \Phi^\top}_{\mathbf{G}} \mathbf{f} \\ &= \Phi \hat{g}(\Lambda) \Phi^\top \mathbf{f} = \hat{g}(\Phi \Lambda \Phi^\top) \mathbf{f} = \hat{g}(\Delta) \mathbf{f} \end{aligned}$$

- Not shift-invariant! ( $\mathbf{G}$  has no circulant structure)
- Commutes with Laplacian:  $\mathbf{G} \Delta \mathbf{f} = \Delta \mathbf{G} \mathbf{f}$
- Filter coefficients depend on basis  $\phi_1, \dots, \phi_n$
- Expensive computation (no FFT):  $O(n^2)$

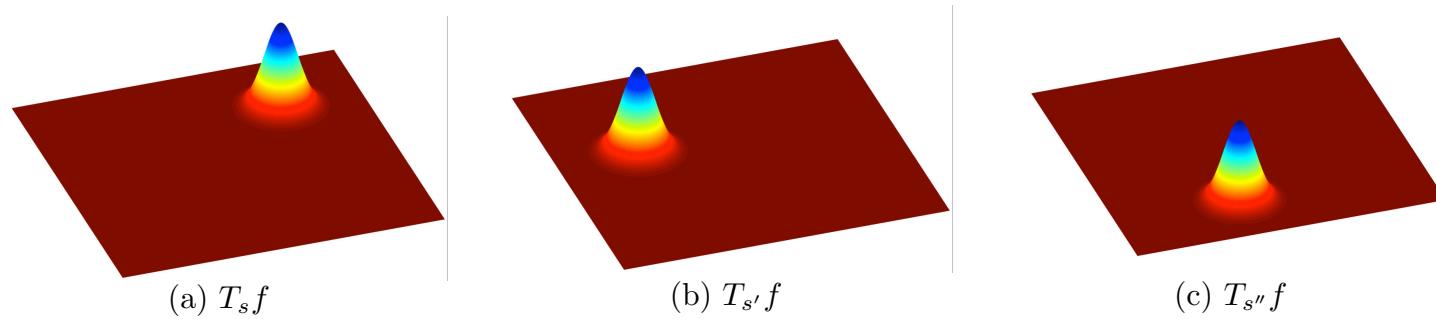
[11] Hammond, Vandergheynst, Gribonval, 2011

# No shift invariance on graphs

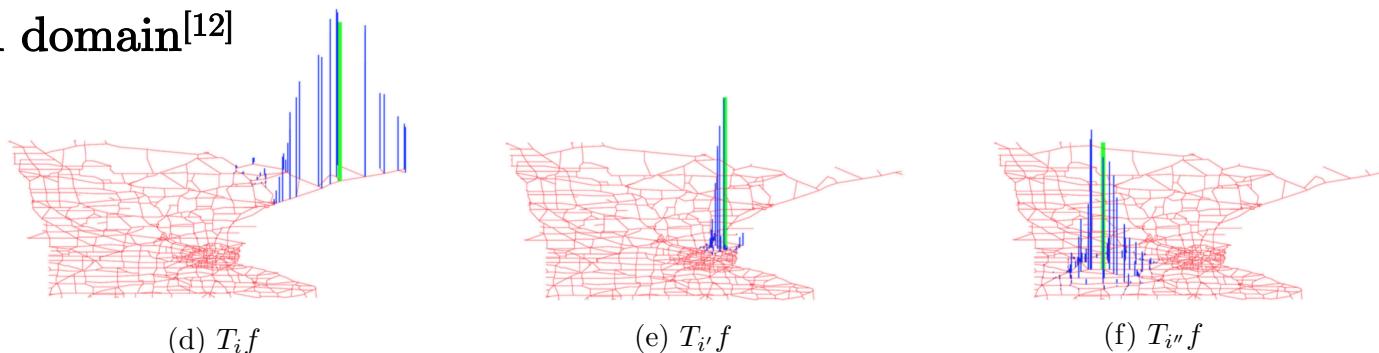
- A signal  $f$  on graph can be translated to vertex  $i$ :

$$T_i f = f \star \delta_i$$

- Euclidean domain



- Graph domain<sup>[12]</sup>



[12] Shuman et al. 2016

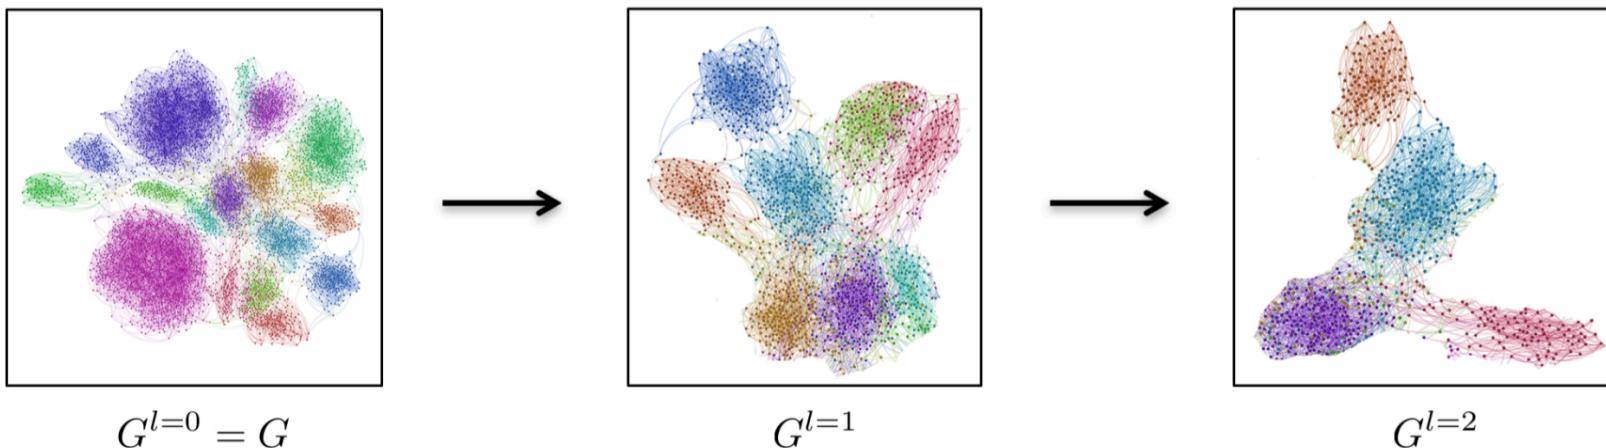
# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- **Spectral Graph Theory**
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- Spectral ConvNets
  - *SplineNets*
  - *ChebNets\* [NIPS'16]*
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\* [NIPS'17]*
- Conclusion

# Graph coarsening for graph pooling

- Goals:
  - Pool similar local features (**max pooling**).
  - Series of pooling layers create **invariance to global geometric deformations**.
- Challenges:
  - Design a **multi-scale coarsening** algorithm that preserves **non-linear** graph structures.
  - How to make graph pooling **fast**?

# Graph coarsening = graph partitioning



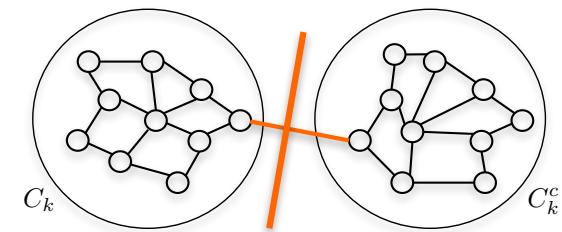
- Graph partitioning decomposes  $G$  into smaller meaningful clusters.
- NP-hard  $\Rightarrow$  Approximation

# Balanced cuts

- Powerful combinatorial graph partitioning models:

- Normalized cut<sup>[13]</sup>

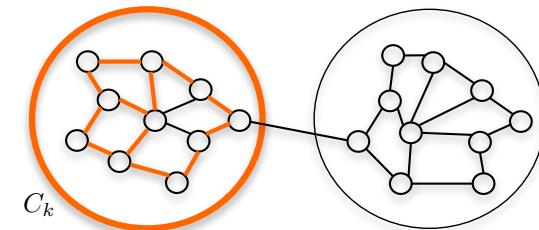
$$\min_{C_1, \dots, C_K} \sum_{k=1}^K \frac{\text{Cut}(C_k, C_k^c)}{\text{Vol}(C_k)}$$



Partitioning by min edge cuts.

- Normalized association

$$\max_{C_1, \dots, C_K} \sum_{k=1}^K \frac{\text{Assoc}(C_k)}{\text{Vol}(C_k)}$$



Partitioning by max vertex matching.

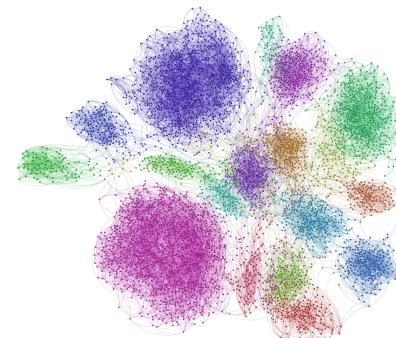
where  $\text{Cut}(A, B) := \sum_{i \in A, j \in B} a_{ij}$ ,  $\text{Assoc}(A) := \sum_{i \in A, i \in B} a_{ij}$ ,  
 $\text{Vol}(A) := \sum_{i \in A, j \in B} d_i$ , and  $d_i := \sum_{j \in V} a_{ij}$ .

- Both models are equivalent, but lead to different algorithms.

[13] Shi, Malik, 2000

# Balanced cuts

- Balanced cuts are NP-hard  $\Rightarrow$  most popular approximation techniques focus on **linear spectral relaxation** (eigenproblem with global solution).
- Graph geometry are generally **not linear**  $\Rightarrow$  **Graclus<sup>[14]</sup>** algorithm computes non-linear clusters that locally maximize the Normalized Association.



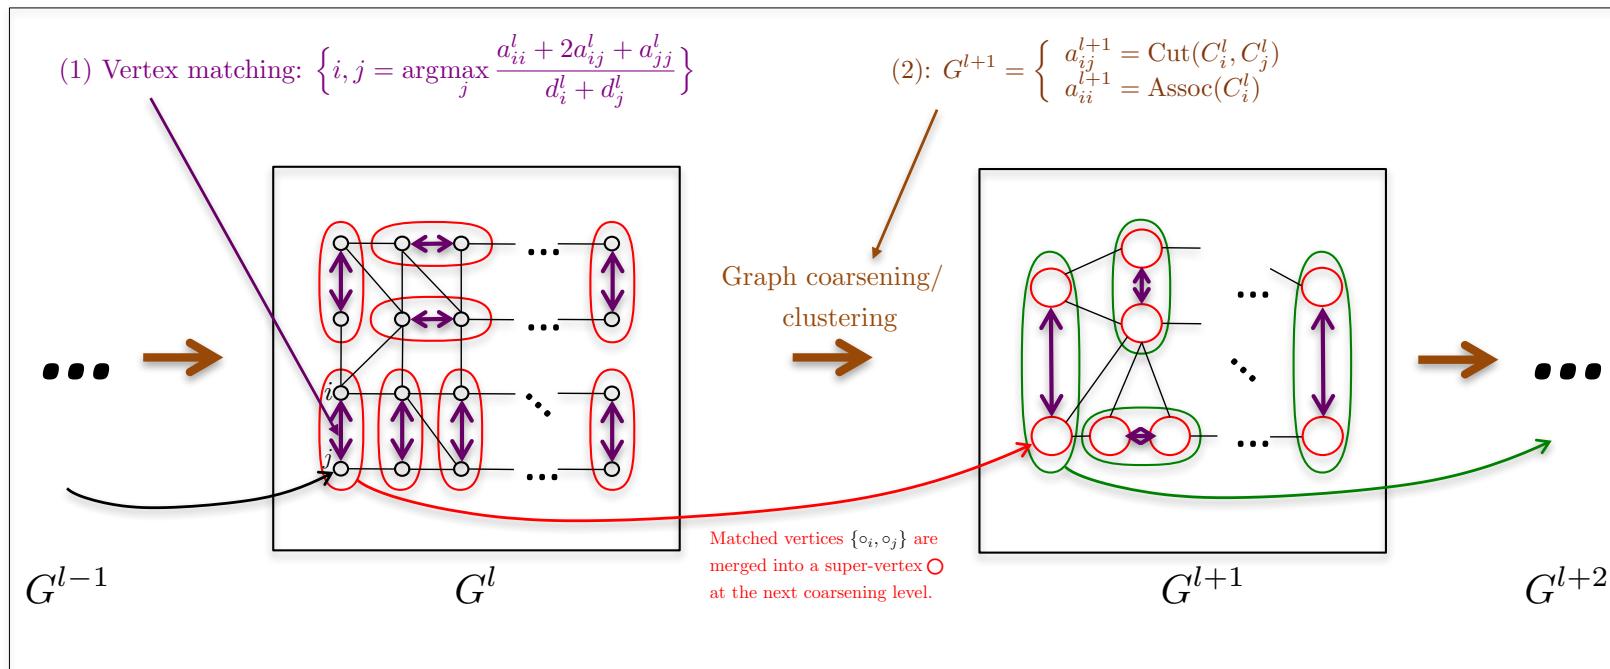
- Graclus algorithm offers a control of the **coarsening ratio of  $\approx 2$**  (like image grid) using **heavy-edge matching<sup>[15]</sup>**.

[14] Dhillon, Guan, Kulis 2007

[15] Karypis, Kumar 1995

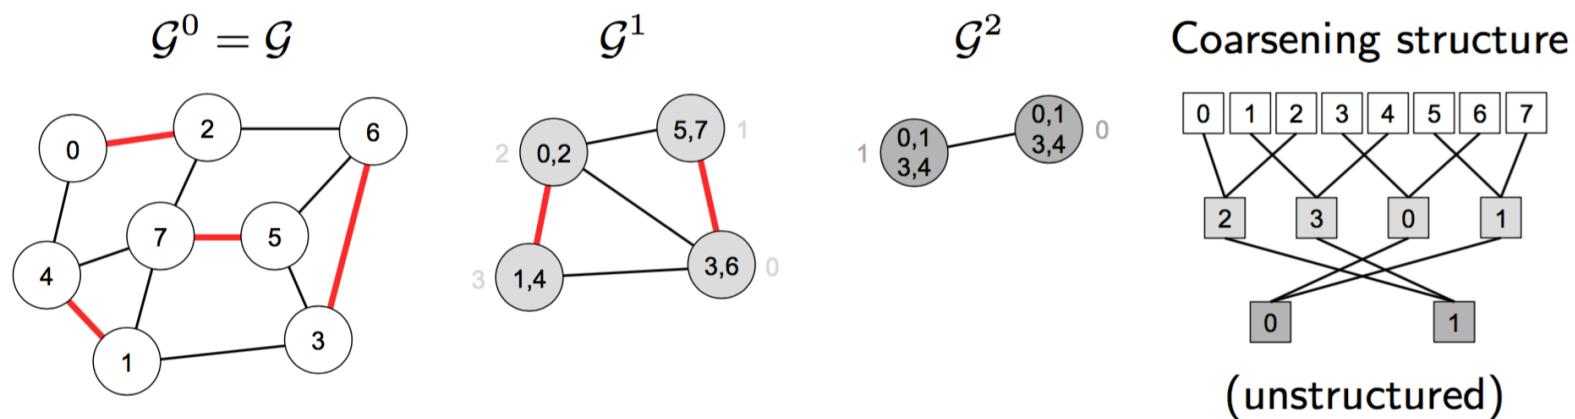
# Heavy-Edge Matching (HEM)

- HEM proceeds by two successive steps, vertex matching and graph coarsening (that guarantees a local solution of Norm assoc):



# Unstructured pooling

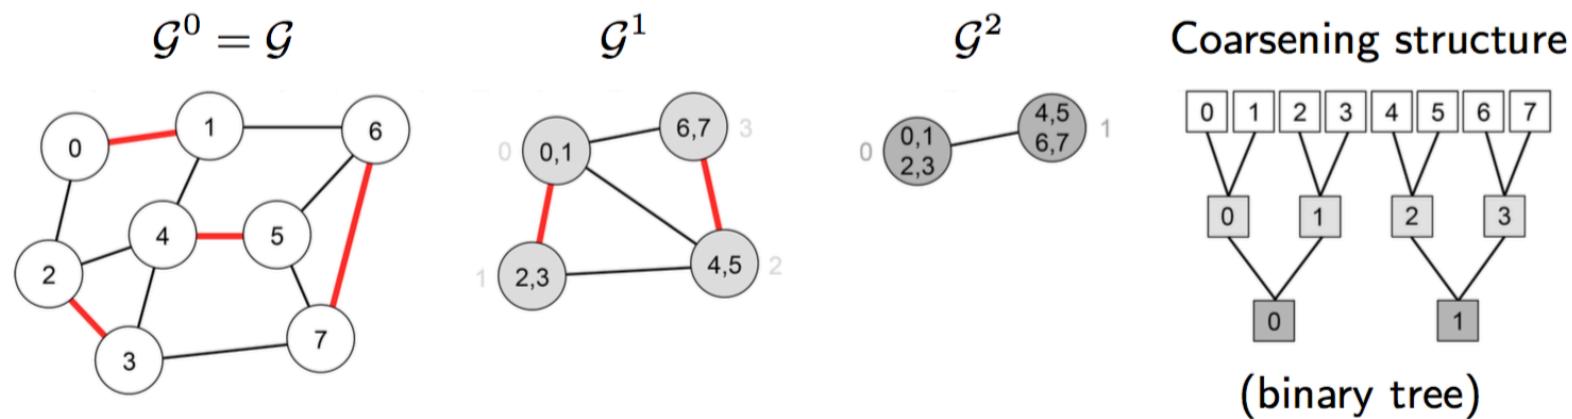
- Sequence of coarsened graphs produced by HEM:



- Stores a table of indices for graph and all its coarsened versions
- Computationally inefficient

# Fast graph pooling<sup>[18]</sup>

- **Structured pooling:** Arrangement of the node indexing such that adjacent nodes are hierarchically merged at the next coarser level.



☺ As efficient as 1D-Euclidean grid pooling.

[18] Defferrard, Bresson, Vandergheynst 2016

# Outline

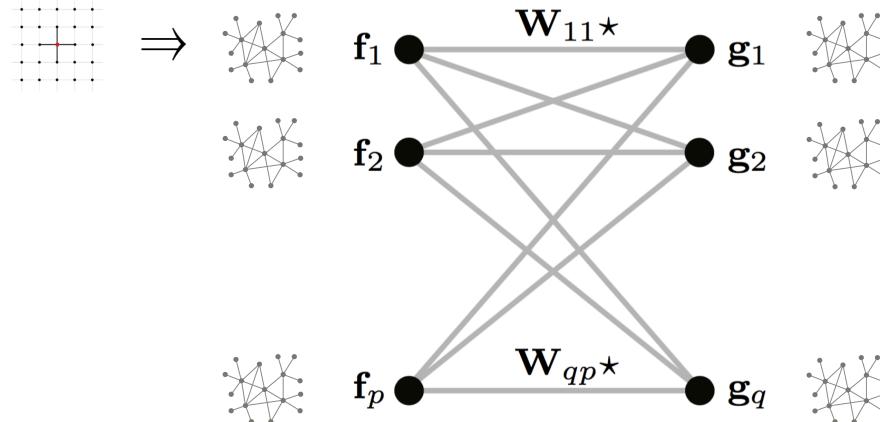
- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- Spectral Graph Theory
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- **Spectral ConvNets**
  - *SplineNets*
  - *ChebNets\* [NIPS'16]*
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\* [NIPS'17]*
- Conclusion

# Vanilla spectral graph ConvNets<sup>[16]</sup>

- Graph convolutional layer

$\mathbf{f}_l$  =  $l$ -th data feature on graphs,  $\dim(\mathbf{f}_l) = n \times 1$

$\mathbf{g}_l$  =  $l$ -th feature map,  $\dim(\mathbf{g}_l) = n \times 1$



Conv. layer

$$\mathbf{g}_l = \xi \left( \sum_{l'=1}^p \mathbf{W}_{l,l'} \star \mathbf{f}_{l'} \right)$$

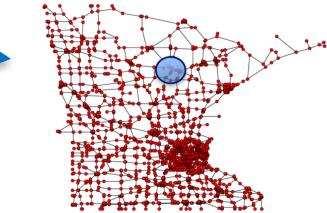
Activation, e.g.  $\xi(x) = \max\{x, 0\}$  rectified linear unit (ReLU)

[16] Bruna, Zaremba, Szlam, LeCun 2014

# Spectral graph convolution

- Convolutional layer in the spatial domain:

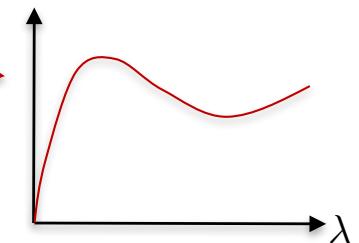
$$\mathbf{g}_l = \xi \left( \sum_{l'=1}^p \mathbf{W}_{l,l'} \star \mathbf{f}_{l'} \right),$$



where  $\mathbf{W}_{l,l'}$  = matrix of graph **spatial filter**,

can also be expressed in the spectral domain (using  $\mathbf{g} \star \mathbf{f} = \Phi \hat{\mathbf{g}}(\Lambda) \Phi^\top \mathbf{f}$ )

$$\mathbf{g}_l = \xi \left( \sum_{l'=1}^p \Phi \hat{\mathbf{W}}_{l,l'} \Phi^\top \mathbf{f}_{l'} \right),$$



where  $\hat{\mathbf{W}}_{l,l'} = n \times n$  diagonal matrix of graph **spectral filter**.

We will denote the spectral filter without the hat symbol, i.e.  $\mathbf{W}_{l,l'}$

# Vanilla spectral graph ConvNets<sup>[16]</sup>

- Series of spectral convolutional layers

$$\mathbf{g}_l^{(k)} = \xi \left( \sum_{l'=1}^{q^{(k-1)}} \Phi \mathbf{W}_{l,l'}^{(k)} \Phi^\top \mathbf{g}_{l'}^{(k-1)} \right),$$

with spectral coefficients  $\mathbf{W}_{l,l'}^{(k)}$  to be learned at each layer.

- ☺ First spectral graph CNN architecture
- ☹ No guarantee of spatial localization of filters
- ☹  $O(n)$  parameters per layer
- ☹  $O(n^2)$  computation of forward and inverse Fourier transforms  $\Phi, \Phi^\top$  (no FFT on graphs)
- ☹ Filters are basis-dependent  $\Rightarrow$  does not generalize across graphs

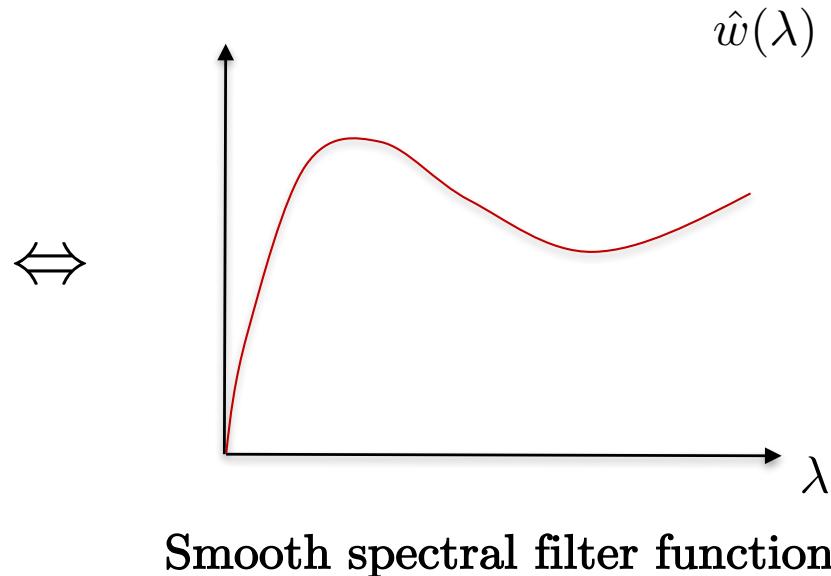
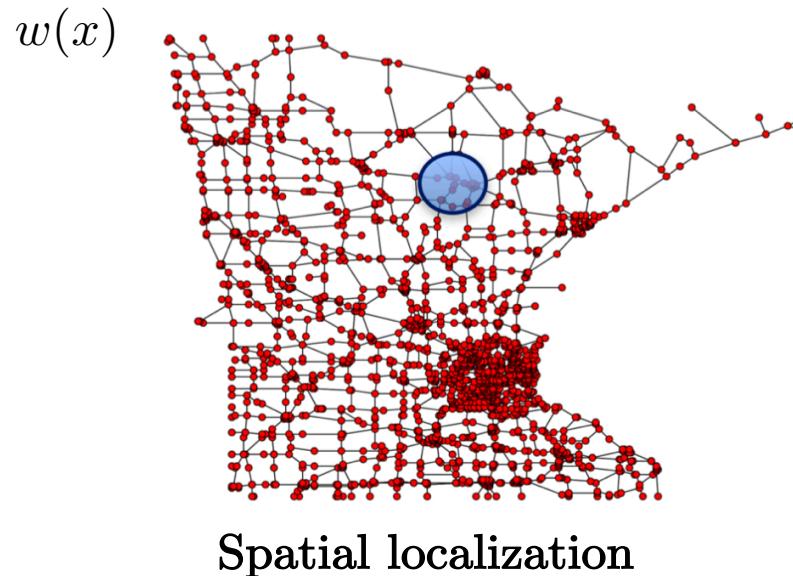
[16] Bruna, Zaremba, Szlam, LeCun 2014

# Spatial localization and spectral smoothness<sup>[17]</sup>

- In the Euclidean setting (by Parseval's identity)

$$\int_{-\infty}^{+\infty} |x|^{2k} |w(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{w}(\lambda)}{\partial \lambda^k} \right|^2 d\lambda$$

⇒ Localization in space = smoothness in frequency domain



[17] Henaff, Bruna, LeCun 2015

# Smooth parametric spectral filter

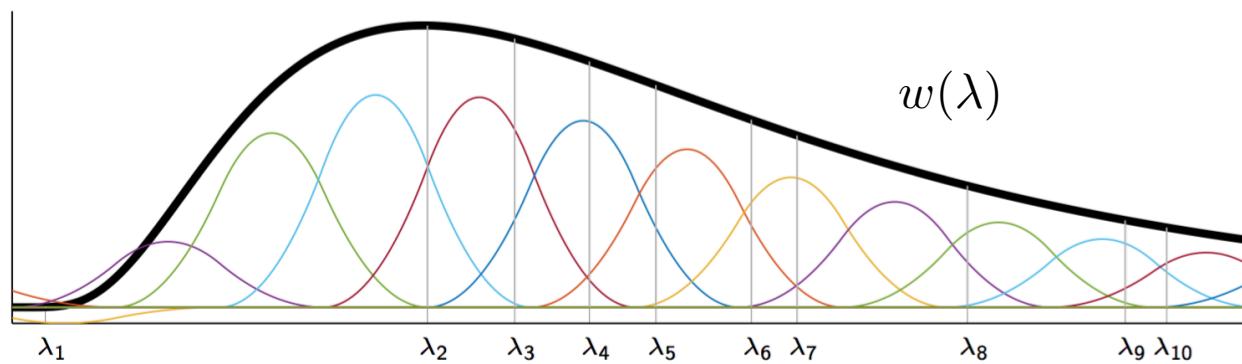
- Parametrize the smooth spectral filter function  $w(\lambda)$  with a linear combination of smooth kernel functions  $\beta_1(\lambda), \dots, \beta_r(\lambda)$ , e.g. splines<sup>[17]</sup>

$$w_{\alpha}(\lambda) = \sum_{j=1}^r \alpha_j \beta_j(\lambda)$$

$\Downarrow$

$$w_{\alpha}(\lambda_i) = \sum_{j=1}^r \alpha_j \beta_j(\lambda_i) = (\mathbf{B}\alpha)_i \quad \Rightarrow \quad \mathbf{W} = \text{Diag}(\mathbf{B}\alpha)$$

where  $\alpha = (\alpha_1, \dots, \alpha_r)^{\top}$  is the vector of filter parameters



[17] (Litman, Bronstein, 2014); Henaff, Bruna, LeCun 2015

# Smooth parametric spectral filter

- **Graph convolutional layer:** Apply the spectral filter  $w$  to a feature signal  $\mathbf{f}$

$$\begin{aligned} w(\Delta)\mathbf{f} &= \Phi w(\Lambda) \Phi^\top \mathbf{f} \\ &= \Phi \begin{pmatrix} w(\lambda_1) & & \\ & \ddots & \\ & & w(\lambda_n) \end{pmatrix} \Phi^\top \mathbf{f} \end{aligned}$$

with  $\alpha$  learnable parameters

$$w_\alpha(\Delta)\mathbf{f} = \Phi \begin{pmatrix} w_\alpha(\lambda_1) & & \\ & \ddots & \\ & & w_\alpha(\lambda_n) \end{pmatrix} \Phi^\top \mathbf{f}$$

# SplineNets<sup>[17]</sup>

- Series of spectral convolutional layers

$$\mathbf{g}_l^{(k)} = \xi \left( \sum_{l'=1}^{q^{(k-1)}} \Phi \mathbf{W}_{l,l'}^{(k)} \Phi^\top \mathbf{g}_{l'}^{(k-1)} \right),$$

with smooth spectral parametric coefficients  $\mathbf{W}_{l,l'}^{(k)}$  to be learned at each layer.

- ☺ Fast-decaying filters in space
- ☺ O(1) parameters per layer
- ☺ O( $n^2$ ) computation of forward and inverse Fourier transforms  
 $\phi, \phi^\top$  (no FFT on graphs)
- ☺ Filters are basis-dependent  $\Rightarrow$  does not generalize across graphs

[17] Henaff, Bruna, LeCun 2015

# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- Spectral Graph Theory
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- **Spectral ConvNets**
  - *SplineNets*
  - *ChebNets\** [NIPS'16]
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\** [NIPS'17]
- Conclusion

# Spectral polynomial filters<sup>[18]</sup>

- Represent smooth spectral functions with polynomials of Laplacian eigenvalues

$$w_{\alpha}(\lambda) = \sum_{j=0}^r \alpha_j \lambda^j$$

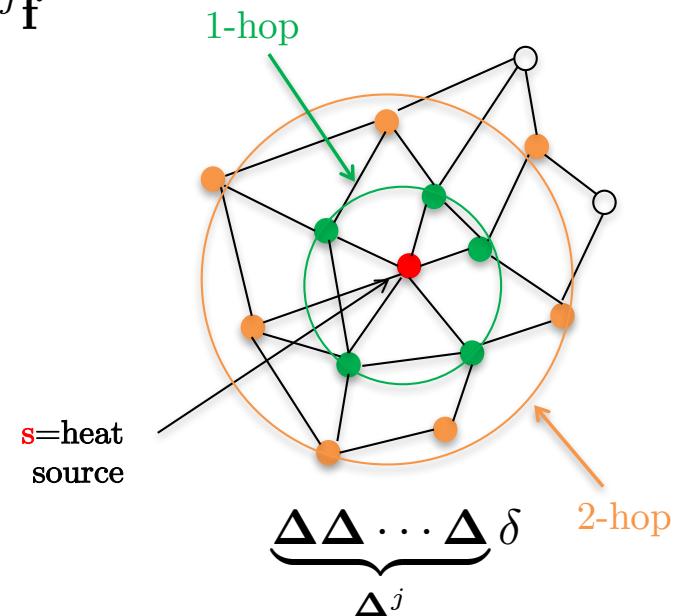
where  $\alpha = (\alpha_1, \dots, \alpha_r)^\top$  is the vector of filter parameters.

- Convolutional layer: Apply spectral filter to feature signal  $f$

$$w_{\alpha}(\Delta)f = \sum_{j=0}^r \alpha_j \Delta^j f$$

- Key observation: Each Laplacian operation increases the support of a function by 1-hop  $\Rightarrow$  Exact control the size of Laplacian-based filters.

[18] Defferrard, Bresson, Vandergheynst 2016



# Linear complexity

- Application of the filter to a feature signal  $f$

$$w_{\alpha}(\Delta)f = \sum_{j=0}^r \alpha_j \Delta^j f$$

- Denote  $X_0 = f$  and define  $X_1 = \Delta X_0 = \Delta f$  and the sequence  $X_j = \Delta X_{j-1}$

$$w_{\alpha}(\Delta)f = \sum_{j=0}^r \alpha_j X_j$$

- Two important observations:
  1. \*No\* need to compute the eigendecomposition of the Laplacian  $(\phi, \Lambda)$ .
  2. Observe that  $\{X_j\}$  are generated by **multiplication of a sparse matrix and a vector**  $\Rightarrow$  Complexity is  $O(Er) = O(n)$  for sparse graphs.
- Graph convolutional layers are **GPU friendly**.

# Spectral graph ConvNets with polynomial filters

- Series of spectral convolutional layers

$$\mathbf{g}_l^{(k)} = \xi \left( \sum_{l'=1}^{q^{(k-1)}} \Phi \mathbf{W}_{l,l'}^{(k)} \Phi^\top \mathbf{g}_{l'}^{(k-1)} \right),$$

with spectral **polynomial** coefficients  $\mathbf{W}_{l,l'}^{(k)}$  to be learned at each layer.

- ☺ Filters are exactly localized in  $r$ -hops support
- ☺  $O(1)$  parameters per layer
- ☺ No computation of  $\phi, \phi^\top \Rightarrow O(n)$  computational complexity  
(assuming sparsely-connected graphs)
- ☹ Unstable under coefficients perturbation (hard to optimize)
- ☹ Filters are basis-dependent  $\Rightarrow$  does not generalize across graphs

[18] Defferrard, Bresson, Vandergheynst 2016

# Chebyshev polynomials

- Graph convolution with (non-orthogonal) **monomial** basis  $1, x, x^2, x^3, \dots$

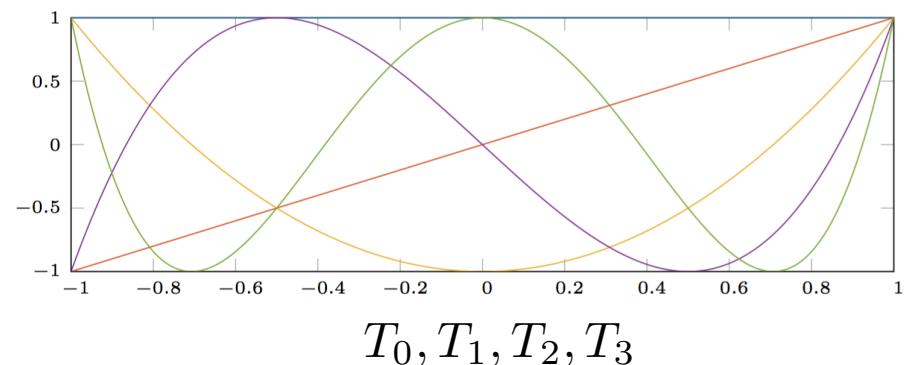
$$w_{\alpha}(\Delta)\mathbf{f} = \sum_{j=0}^r \alpha_j \Delta^j \mathbf{f}$$

- Graph convolution with (orthogonal) **Chebyshev polynomials**

$$w_{\alpha}(\tilde{\Delta})\mathbf{f} = \sum_{j=0}^r \alpha_j T_j(\tilde{\Delta})\mathbf{f}$$

- Orthonormal on  $L^2([-1, +1])$  w.r.t.  $\langle f, g \rangle = \int_{-1}^{+1} f(\tilde{\lambda})g(\tilde{\lambda}) \frac{d\tilde{\lambda}}{\sqrt{1-\tilde{\lambda}^2}}$

- Stable under perturbation of coefficients



# ChebNets<sup>[18]</sup>

- Application of the filter with the scaled Laplacian  $\tilde{\Delta} = 2\lambda_n^{-1}\Delta - \mathbf{I}$

$$w_{\alpha}(\tilde{\Delta})\mathbf{f} = \sum_{j=0}^r \alpha_j T_j(\tilde{\Delta})\mathbf{f} = \sum_{j=0}^r \alpha_j \mathbf{X}^{(j)}$$

with

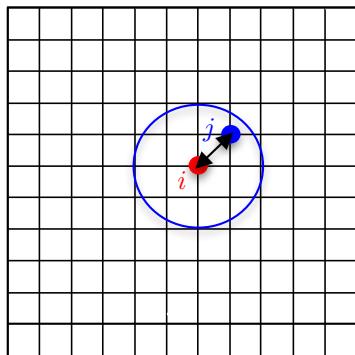
$$\begin{aligned}\mathbf{X}^{(j)} &= T_j(\tilde{\Delta})\mathbf{f} \\ &= 2\tilde{\Delta}\mathbf{X}^{(j-1)} - \mathbf{X}^{(j-2)}, \quad \mathbf{X}^{(0)} = \mathbf{f}, \quad \mathbf{X}^{(1)} = \tilde{\Delta}\mathbf{f}\end{aligned}$$

- ☺ Filters are exactly localized in  $r$ -hops support
- ☺  $O(1)$  parameters per layer
- ☺ No computation of  $\phi, \phi^\top \Rightarrow O(n)$  computational complexity  
(assuming sparsely-connected graphs)
- ☺ Stable under coefficients perturbation
- ☹ Filters are basis-dependent  $\Rightarrow$  does not generalize across graphs

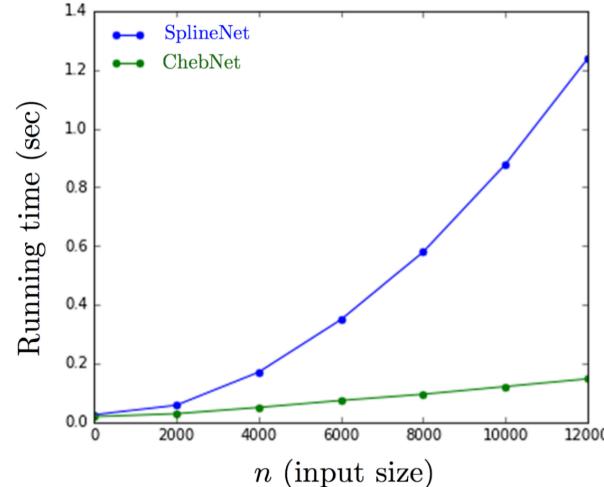
[18] Defferrard, Bresson, Vandergheynst 2016

# Numerical experiments

Graph: a 8-NN graph of the Euclidean grid



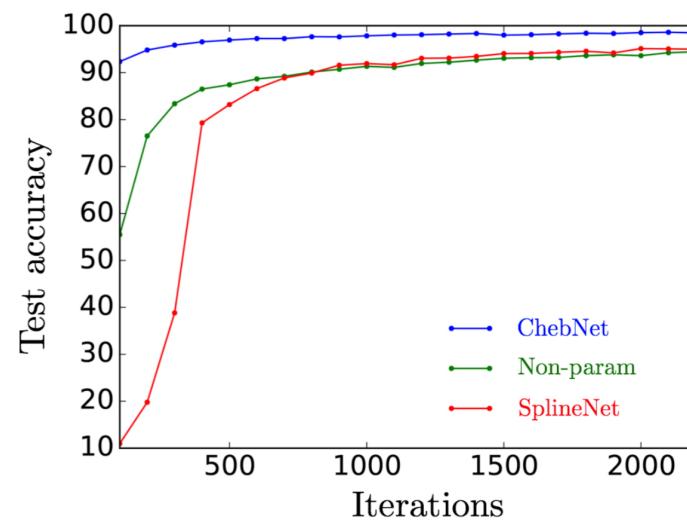
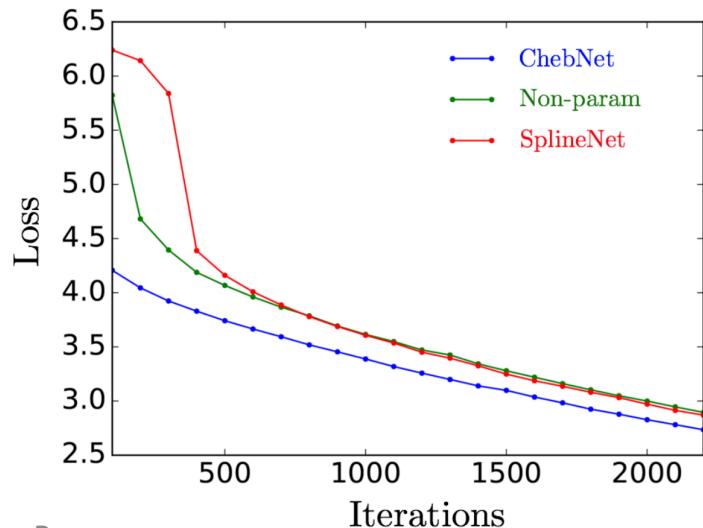
- Running time



- Accuracy

| Model     | Order | Accuracy |
|-----------|-------|----------|
| LeNet5    | -     | 99.33%   |
| SplineNet | 25    | 97.75%   |
| ChebNet   | 25    | 99.14%   |

- Optimization



# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- Spectral Graph Theory
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- **Spectral ConvNets**
  - *SplineNets*
  - *ChebNets\** [NIPS'16]
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\** [NIPS'17]
- Conclusion

## Graph convolutional nets<sup>[19]</sup>: simplified ChebNets

- Use Chebychev polynomials of degree  $r=2$  and assume  $\lambda_n \approx 2$

$$\begin{aligned} w_{\alpha}(\Delta)\mathbf{f} &= \alpha_0\mathbf{f} + \alpha_1(\Delta - \mathbf{I})\mathbf{f} \\ &= \alpha_0\mathbf{f} - \alpha_1\mathbf{D}^{-1/2}\mathbf{W}\mathbf{D}^{-1/2}\mathbf{f} \end{aligned}$$

- Further constrain  $\alpha = \alpha_0 = -\alpha_1$  to obtain a single-parameter filter

$$w_{\alpha}(\Delta)\mathbf{f} = \alpha(\mathbf{I} + \mathbf{D}^{-1/2}\mathbf{W}\mathbf{D}^{-1/2})\mathbf{f}$$

- Caveat: The eigenvalues of  $\mathbf{I} + \mathbf{D}^{-1/2}\mathbf{W}\mathbf{D}^{-1/2}$  are now in  $[0,2]$   
⇒ repeated application of the filter results in **numerical instability**
- Fix: Apply a **renormalization**

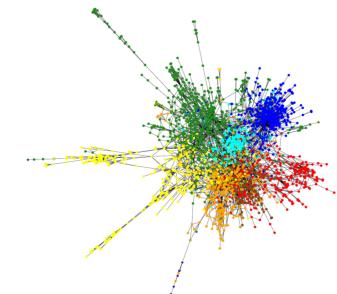
$$w_{\alpha}(\Delta)\mathbf{f} = \alpha\tilde{\mathbf{D}}^{-1/2}\tilde{\mathbf{W}}\tilde{\mathbf{D}}^{-1/2}\mathbf{f}$$

with  $\tilde{\mathbf{W}} = \mathbf{W} + \mathbf{I}$  and  $\tilde{\mathbf{D}} = \text{diag}(\sum_{j \neq i} \tilde{w}_{ij})$

[19] Kipf, Welling 2016

# Example: citation networks

| Method                                     | Cora <sup>1</sup> | PubMed <sup>2</sup> |
|--|-------------------|---------------------|
| Manifold Regularization <sup>3</sup>       | 59.5%             | 70.7%               |
| Semidefinite Embedding <sup>4</sup>        | 59.0%             | 71.1%               |
| Label Propagation <sup>5</sup>             | 68.0%             | 63.0%               |
| DeepWalk <sup>6</sup>                      | 67.2%             | 65.3%               |
| Planetoid <sup>7</sup>                     | 75.7%             | 77.2%               |
| <b>Graph Convolutional Net<sup>8</sup></b> | <b>81.59%</b>     | <b>78.72%</b>       |



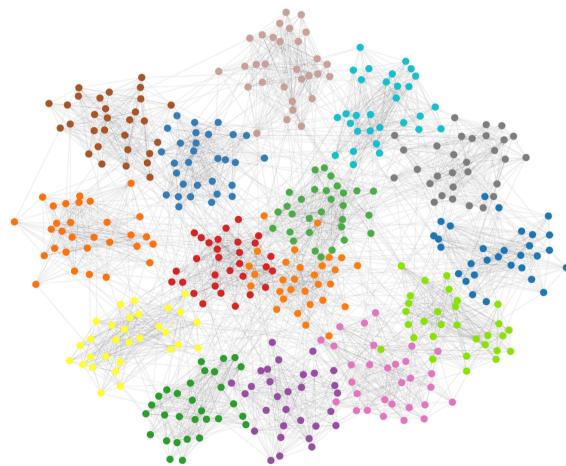
Classification accuracy of different methods on citation network datasets

Monti et al. 2016; data: <sup>1,2</sup>Sen et al. 2008; methods: <sup>3</sup>Belkin et al. 2006; <sup>4</sup>Weston et al. 2012; <sup>5</sup>Zhu et al. 2003; <sup>6</sup>Perozzi et al. 2014; <sup>7</sup>Yang et al. 2016; <sup>8</sup>Kipf, Welling 2016

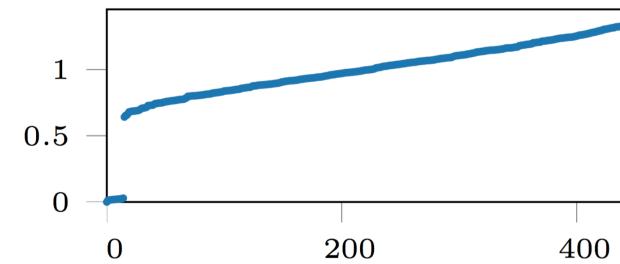
# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- Spectral Graph Theory
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- **Spectral ConvNets**
  - *SplineNets*
  - *ChebNets\** [NIPS'16]
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\** [NIPS'17]
- Conclusion

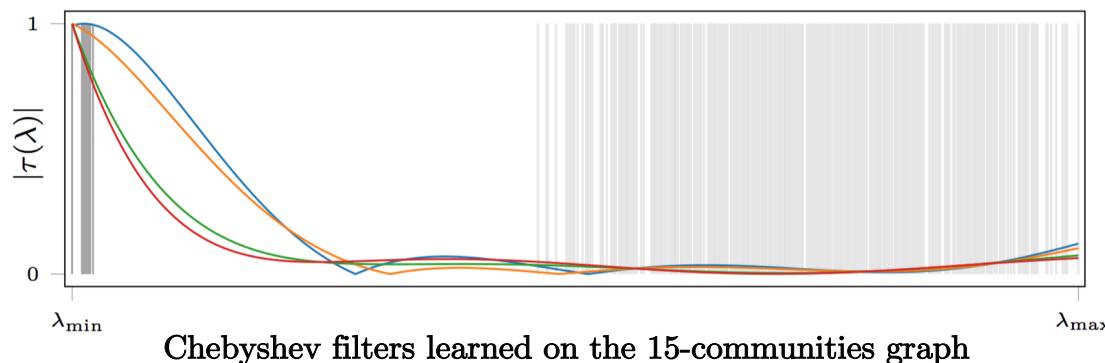
# Community graphs



Synthetic graph with 15 communities



Normalized Laplacian eigenvalues



Chebyshev filters learned on the 15-communities graph

⌚ ChebNet is unstable to produce filters with frequency bands of interest

# Spectral graph ConvNets with Cayley polynomials<sup>[20]</sup>

- Cayley polynomials are a family of real-valued rational functions (w/ complex coefficients)

$$p(\lambda) = c_0 + 2\operatorname{Re}\left\{\sum_{j=1}^r c_j \underbrace{\left(\frac{\lambda - i}{\lambda + i}\right)^j}_{C(\lambda)}\right\}$$

- Cayley transform

$C(\lambda) = \frac{\lambda - i}{\lambda + i}$  is a smooth bijection from  $\mathbb{R}$  to  $e^{i\mathbb{R}} \setminus \{1\}$

- Expressivity: Since  $z^{-1} = \bar{z}$  for  $z \in e^{i\mathbb{R}}$  and  $2\operatorname{Re}\{z\} = z + \bar{z}$  we can rewrite  $p(\lambda)$  as a real trigonometric polynomial w.r.t.  $C(\lambda)$

$$\begin{aligned} p(\lambda) &= c_0 + \sum_{j=1}^r c_j \underbrace{C^j(\lambda)}_{e^{ij\omega_\lambda}} + \bar{c}_j \underbrace{C^{-j}(\lambda)}_{e^{-ij\omega_\lambda}} \\ &= 2 \sum_{j=1}^r \operatorname{Re}\{c_j\} \cos(j\omega_\lambda) - \operatorname{Im}\{c_j\} \sin(j\omega_\lambda) \end{aligned}$$

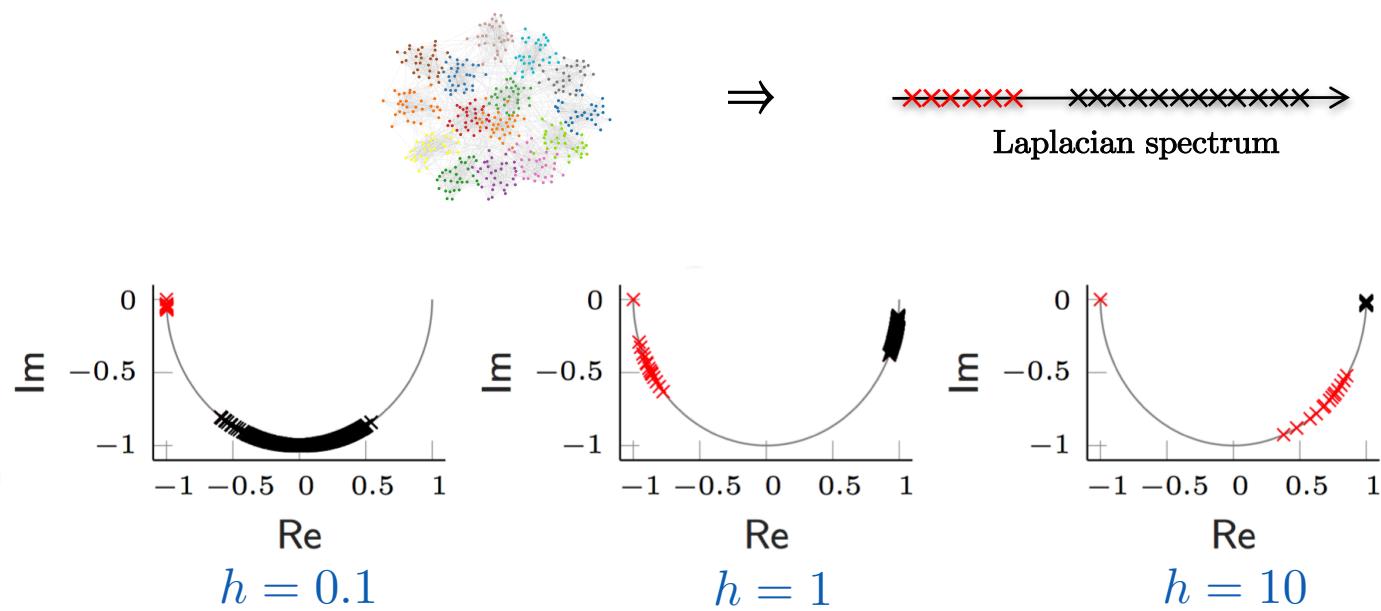
[20] (Cayley 1846); Levie, Monti, Bresson, Bronstein 2017

# Spectral zoom<sup>[20]</sup>

- Applying Cayley transform to the scaled Laplacian  $h\Delta$

$$C(h\Delta) = (h\Delta - i\mathbf{I})(h\Delta + i\mathbf{I})^{-1}$$

results in a non-linear transformation of the eigenvalues  
(spectral zoom)



Cayley transform  $C(h\lambda)$  of the 15-communities graph Laplacian spectrum

[20] Levie, Monti, Bresson, Bronstein 2017

# Fast inversion

- Application of Cayley filter  $w(h\Delta)\mathbf{f}$  requires the solution of recursive linear system

$$\begin{aligned}\mathbf{y}_0 &= \mathbf{f} \\ (h\Delta + i\mathbf{I})\mathbf{y}_j &= (h\Delta - i\mathbf{I})\mathbf{y}_{j-1} \quad j = 1, \dots, r\end{aligned}$$

- Exact (stable) solution:  $O(n^3)$  complexity
- Approximate solution using K Jacobi iterations

$$\begin{aligned}\tilde{\mathbf{y}}_j^{(k+1)} &= \mathbf{J}\tilde{\mathbf{y}}_j^{(k)} + \text{Diag}^{-1}(h\Delta + i\mathbf{I})(h\Delta - i\mathbf{I})\tilde{\mathbf{y}}_{j-1} \\ \tilde{\mathbf{y}}_j^{(0)} &= \text{Diag}^{-1}(h\Delta + i\mathbf{I})(h\Delta - i\mathbf{I})\tilde{\mathbf{y}}_{j-1}\end{aligned}$$

with  $\mathbf{J} = -\text{Diag}^{-1}(h\Delta + i\mathbf{I})\text{Off}(h\Delta + i\mathbf{I})$

- Application of the approximate Cayley filter

$$\sum_{j=0}^r c_j \tilde{\mathbf{y}}_j \approx w(h\Delta)\mathbf{f}$$

Complexity:  $O(ErK) = O(n)$

# CayleyNets<sup>[30]</sup>

- Represent spectral transfer function as a Cayley polynomial of order  $r$

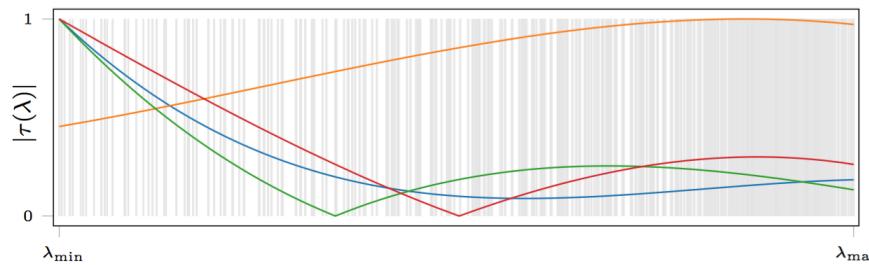
$$w_{\mathbf{c}, h}(\lambda) = c_0 + 2\operatorname{Re} \left\{ \sum_{j=1}^r c_j (h\lambda - i)^j (h\lambda + i)^{-j} \right\}$$

where the filter parameters are the vector of complex coefficients  $\mathbf{c} = (c_0, \dots, c_r)^\top$  and the spectral zoom  $h$ .

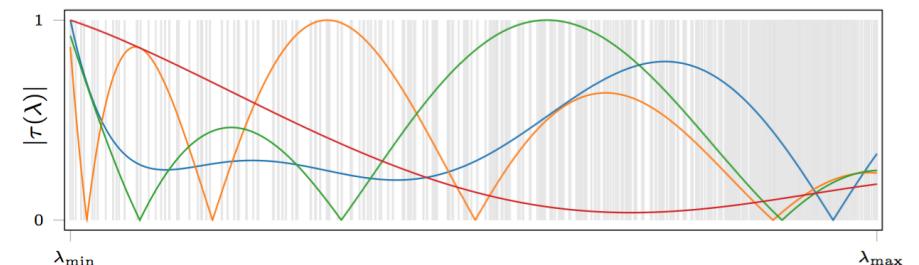
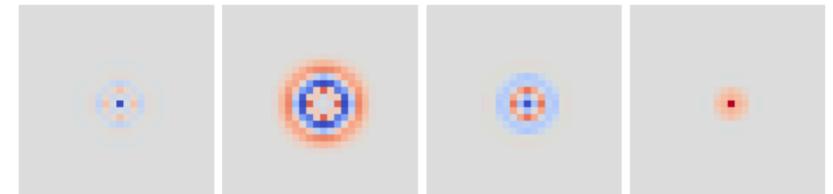
- ☺ Filters have guaranteed exponential spatial decay
- ☺  $O(1)$  parameters per layer
- ☺  $O(n)$  computational complexity with Jacobi approximate inversion (assuming sparsely-connected graph)
- ☺ Spectral zoom property allowing to better localize in frequency
- ☺ Richer class of filters than Chebyshev for the same order
- ☹ Filters are basis-dependent  $\Rightarrow$  does not generalize across graphs

[20] Levie, Monti, Bresson, Bronstein 2017

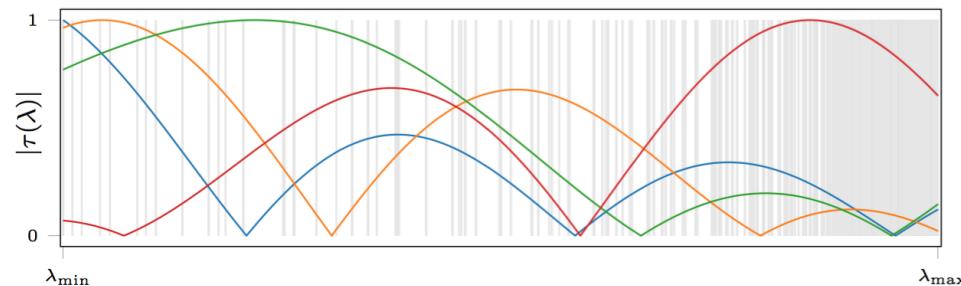
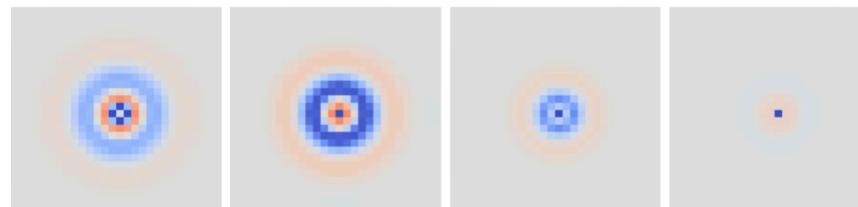
# Chebyshev vs Cayley



Example of Chebyshev filters (order  $r = 3$ ) on Euclidean grid

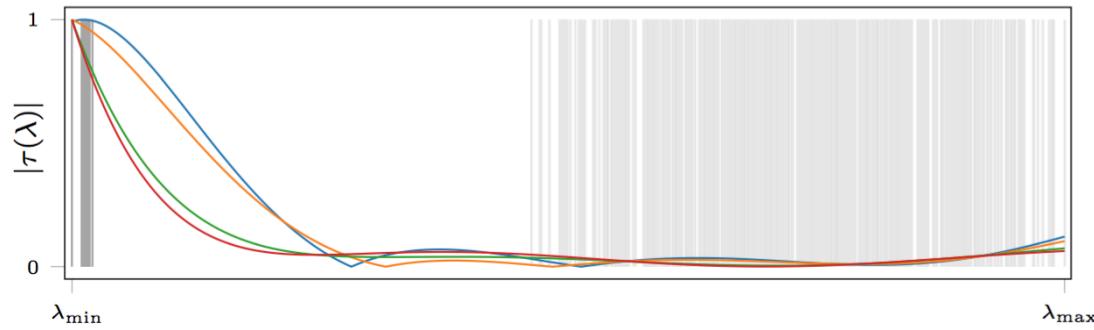


Example of Chebyshev filters (order  $r = 7$ ) on Euclidean grid

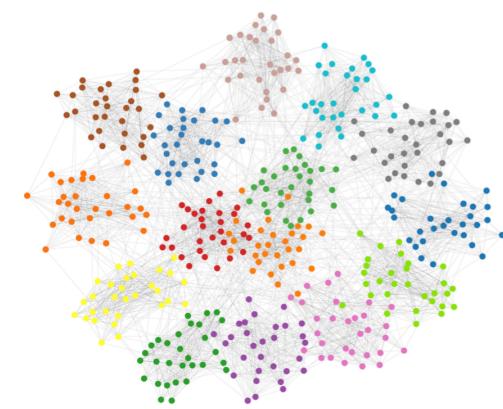


Example of Cayley filters (order  $r = 3$ ) on Euclidean grid

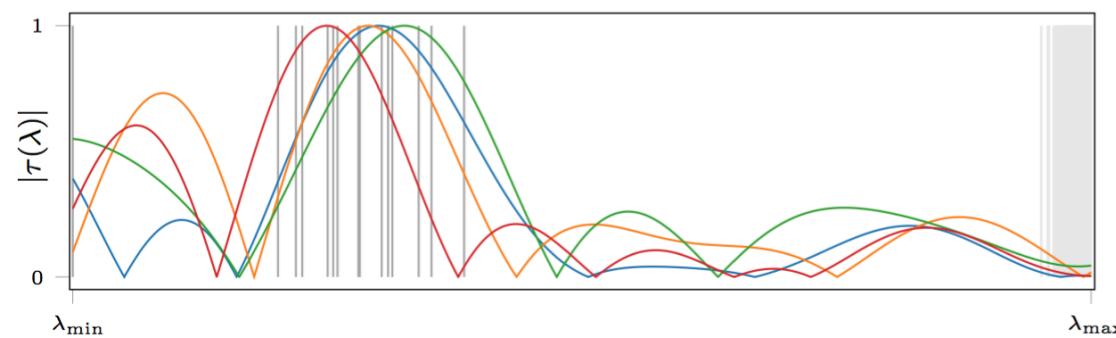
# Chebyshev vs Cayley: community graphs example



Example of Chebyshev filters learned on the 15-communities graph

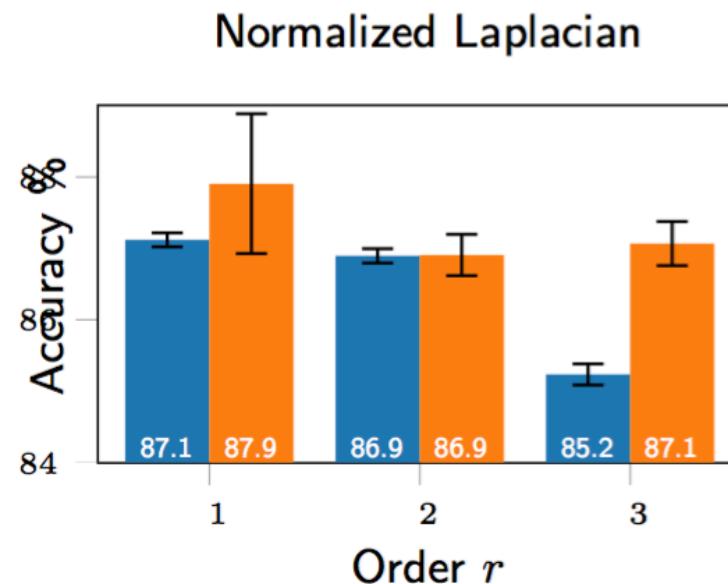


Synthetic graph with 15 communities



Example of Cayley filters learned on the 15-communities graph

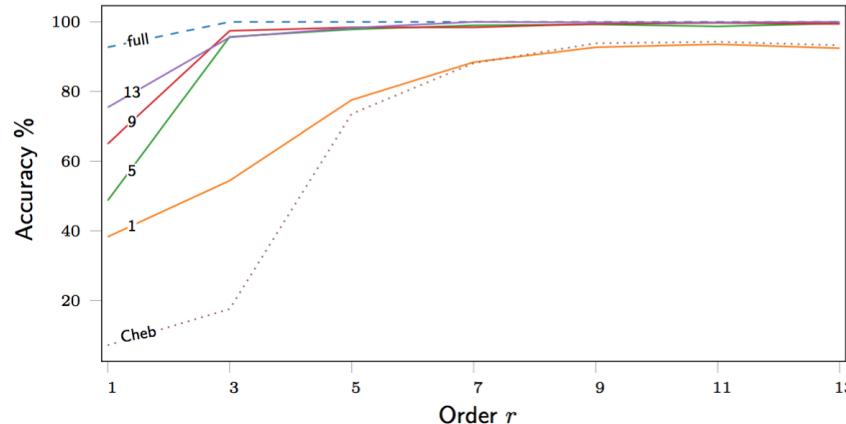
# Chebyshev vs Cayley: Cora example



Accuracy of ChebNet (blue) and CayleyNet (orange)  
on Cora dataset

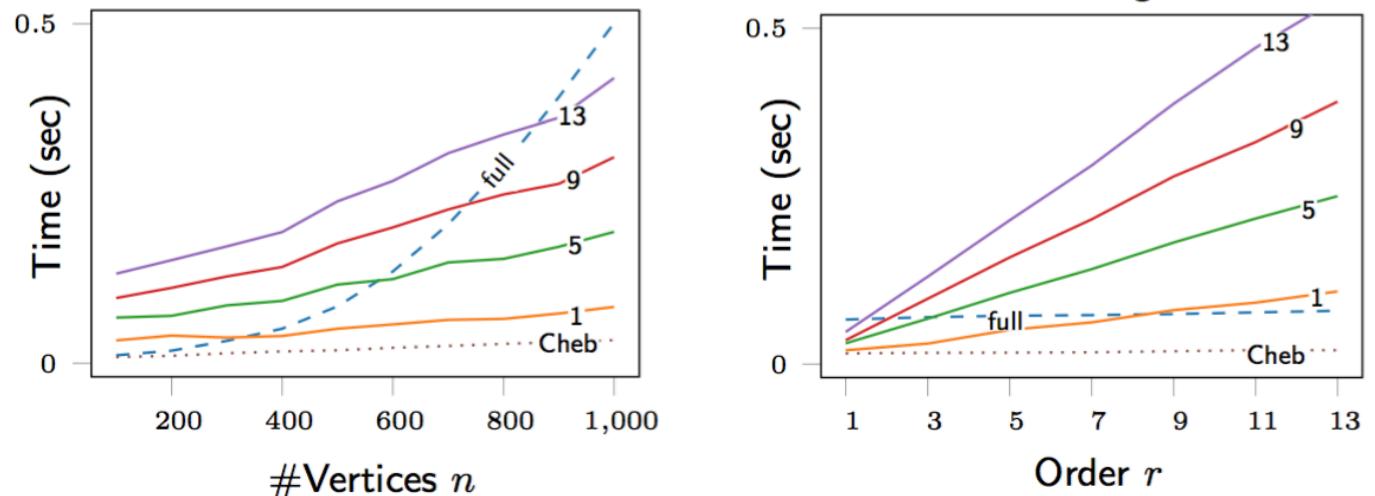
# Approximate inversion

- Accuracy of approximate inversion



Community detection accuracy of CayleyNet using approximate Jacobi inversion on the synthetic 15-community graph

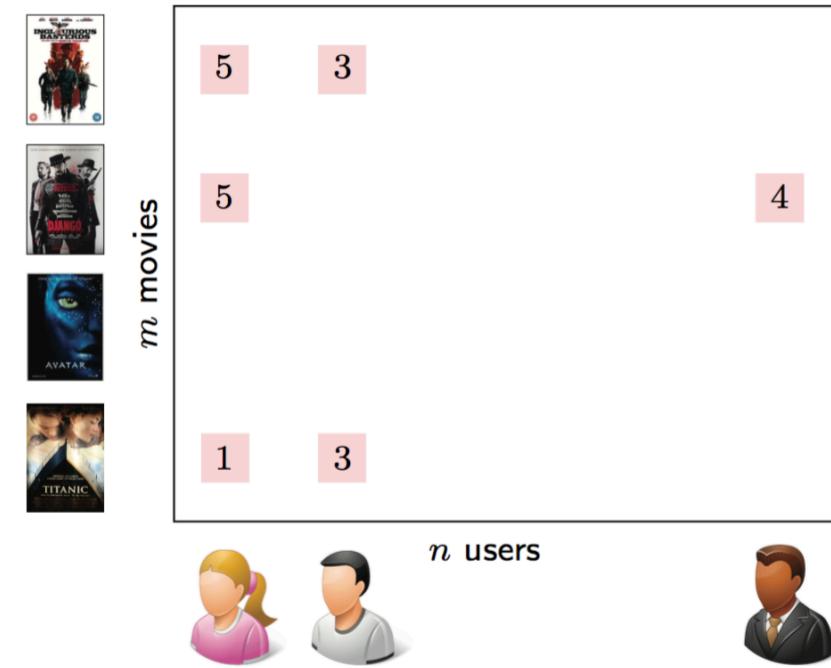
- Speed of approximate inversion



# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- Spectral Graph Theory
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- Spectral ConvNets
  - *SplineNets*
  - *ChebNets\* [NIPS'16]*
  - *GraphConvNets*
  - *CayleyNets\**
- **Spectral ConvNets on Multiple Graphs**
  - *Recommender Systems\* [NIPS'17]*
- Conclusion

# Matrix completion - Netflix challenge



$$\min_{\mathbf{X} \in \mathbb{R}^{m \times n}} \quad \text{rank}(\mathbf{X}) \quad \text{s.t.} \quad x_{ij} = a_{ij} \quad \forall ij \in \Omega$$

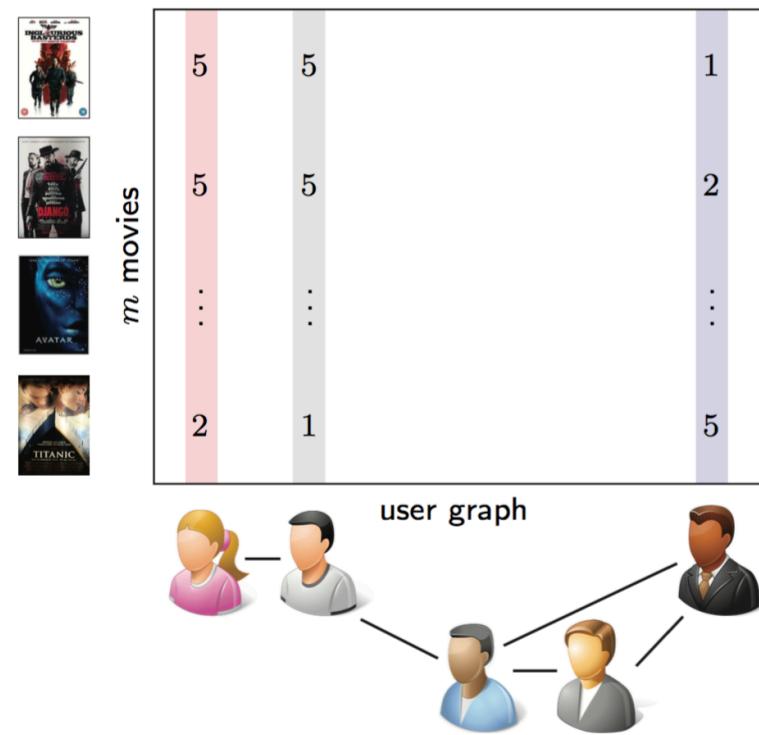
$\downarrow$  Convex relaxation<sup>[21]</sup>

$$\min_{\mathbf{X} \in \mathbb{R}^{m \times n}} \quad \|\mathbf{X}\|_* \quad \text{s.t.} \quad x_{ij} = a_{ij} \quad \forall ij \in \Omega$$

$$\min_{\mathbf{X} \in \mathbb{R}^{m \times n}} \quad \|\mathbf{X}\|_* + \mu \|\mathbf{\Omega} \circ (\mathbf{X} - \mathbf{A})\|_F^2$$

[21] Candes et al. 2008

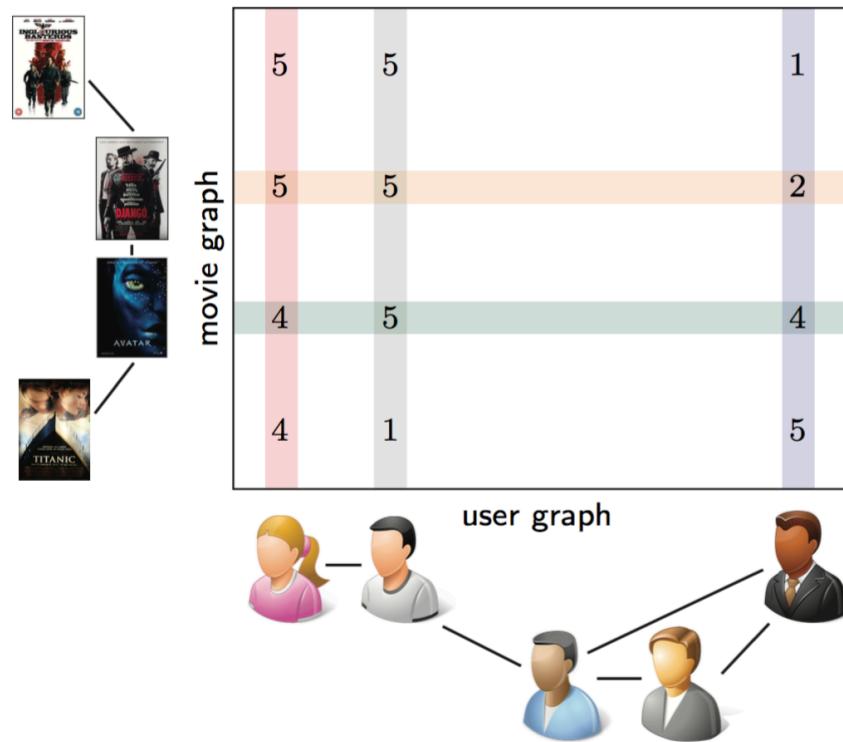
# Matrix completion on single graph<sup>[22]</sup>



$$\min_{\mathbf{X} \in \mathbb{R}^{m \times n}} \quad \|\mathbf{X}\|_* + \mu \|\boldsymbol{\Omega} \circ (\mathbf{X} - \mathbf{A})\|_F^2 + \mu_c \underbrace{\text{tr}(\mathbf{X} \boldsymbol{\Delta}_c \mathbf{X}^\top)}_{\|\mathbf{X}\|_{\mathcal{G}_c}^2}$$

[22] Kalofolias, Bresson, Bronstein, Vandergheynst, 2014

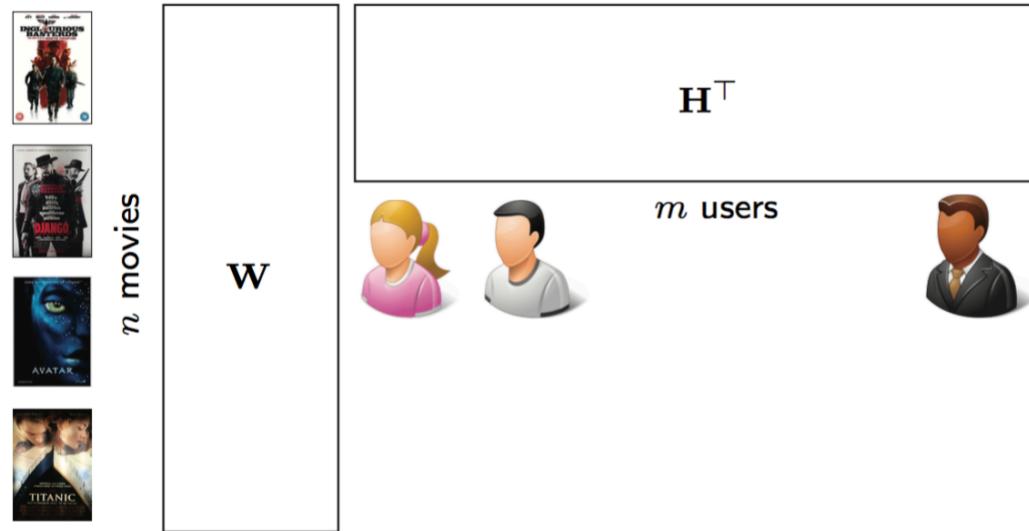
# Matrix completion on multiple graphs<sup>[22]</sup>



$$\min_{\mathbf{X} \in \mathbb{R}^{m \times n}} \|\mathbf{X}\|_* + \mu \|\boldsymbol{\Omega} \circ (\mathbf{X} - \mathbf{A})\|_F^2 + \mu_c \underbrace{\text{tr}(\mathbf{X} \boldsymbol{\Delta}_c \mathbf{X}^\top)}_{\|\mathbf{X}\|_{\mathcal{G}_c}^2} + \mu_r \underbrace{\text{tr}(\mathbf{X}^\top \boldsymbol{\Delta}_r \mathbf{X})}_{\|\mathbf{X}\|_{\mathcal{G}_r}^2}$$

[22] Kalofolias, Bresson, Bronstein, Vandergheynst, 2014

# Factorized matrix completion models<sup>[23]</sup>

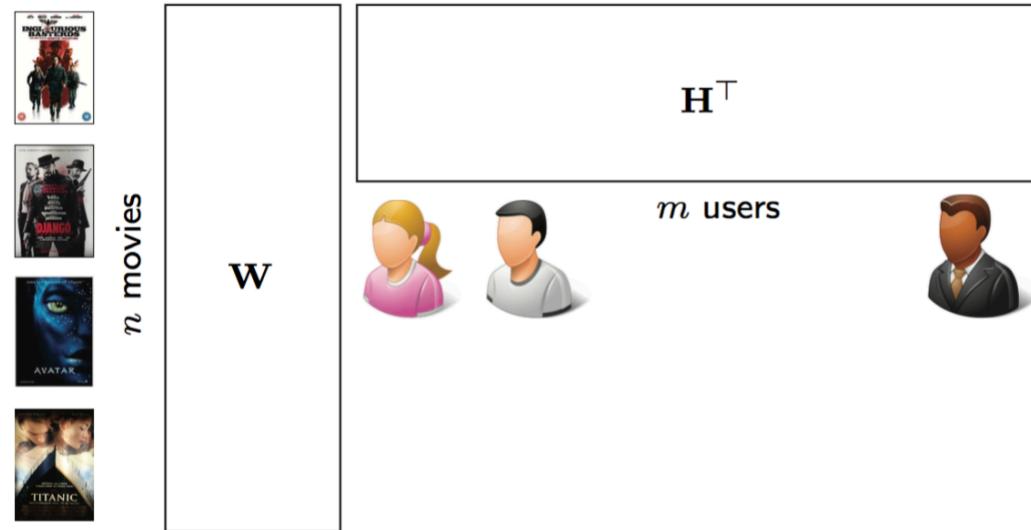


$$\min_{\substack{\mathbf{W} \in \mathbb{R}^{m \times s} \\ \mathbf{H} \in \mathbb{R}^{n \times s}}} \mu \|\Omega \circ (\mathbf{X} - \mathbf{A})\|_F^2 + \mu_c \|\mathbf{W}\|_F^2 + \mu_r \|\mathbf{H}\|_F^2$$

- ☺  $O(n+m)$  variables instead of  $O(nm)$
- ☺  $\text{rank}(\mathbf{X}) = \text{rank}(\mathbf{W}\mathbf{H}^T) \leq s$  by construction

[23] Srebro, Rennie, Jaakkola 2004

# Factorized matrix completion on graphs<sup>[24]</sup>



$$\min_{\substack{\mathbf{W} \in \mathbb{R}^{m \times s} \\ \mathbf{H} \in \mathbb{R}^{n \times s}}} \mu \|\boldsymbol{\Omega} \circ (\mathbf{X} - \mathbf{A})\|_F^2 + \mu_c \text{tr}(\mathbf{H}^\top \boldsymbol{\Delta}_c \mathbf{H}) + \mu_r \text{tr}(\mathbf{W}^\top \boldsymbol{\Delta}_r \mathbf{W})$$

- ☺  $O(n+m)$  variables instead of  $O(nm)$
- ☺  $\text{rank}(\mathbf{X}) = \text{rank}(\mathbf{W}\mathbf{H}^\top) \leq s$  by construction

[24] Rao et al. 2015

# 1D Fourier transform

$$\begin{matrix} \text{Heatmap} \\ (\hat{\mathbf{x}}_1, \dots, \hat{\mathbf{x}}_n) \end{matrix} = \begin{matrix} \Phi^\top \\ \text{Matrix} \end{matrix} \times \begin{matrix} \text{Mona Lisa} \\ \mathbf{x} \end{matrix}$$

Column-wise transform

## 2D Fourier transform

$$\hat{\mathbf{X}} = \Phi^T \times \mathbf{X} \times \Phi$$

The diagram illustrates the 2D Fourier transform process. On the left is a heatmap labeled  $\hat{\mathbf{X}}$ , representing the Fourier transform of the image. To its right is an equals sign. Following the equals sign are three components: a vertical matrix labeled  $\Phi^T$ , a grayscale image of the Mona Lisa labeled  $\mathbf{X}$ , and a horizontal matrix labeled  $\Phi$ . Between the  $\Phi^T$  matrix and the image  $\mathbf{X}$  is a multiplication symbol ( $\times$ ), and between the image  $\mathbf{X}$  and the  $\Phi$  matrix is another multiplication symbol ( $\times$ ). This visualizes the formula  $\hat{\mathbf{X}} = \Phi^T \times \mathbf{X} \times \Phi$ .

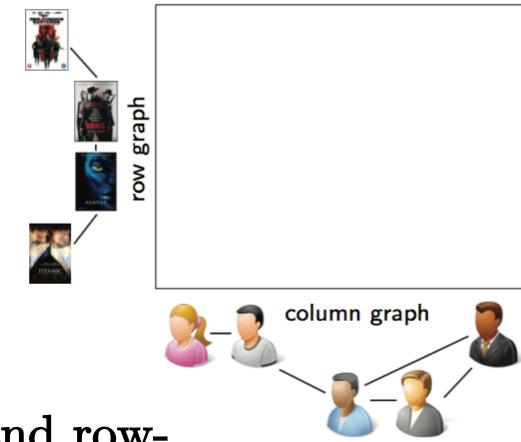
Column-wise transform + Row-wise transform = 2D transform

From 1D signal processing to 2D image processing

# Multi-graph Fourier transform and convolution<sup>[25]</sup>

- Multi-graph Fourier transform

$$\hat{\mathbf{X}} = \Phi_r^\top \mathbf{X} \Phi_c$$



where  $\Phi_c$ ,  $\Phi_r$  are the eigenvectors of the column- and row-graph Laplacians  $\Delta_c$ ,  $\Delta_r$ , respectively.

- Multi-graph spectral convolution

$$\begin{aligned}\mathbf{X} \star \mathbf{Y} &= \Phi_r ((\Phi_r^\top \mathbf{X} \Phi_c) \circ (\Phi_r^\top \mathbf{Y} \Phi_c)) \Phi_c^\top \\ &= \Phi_r (\hat{\mathbf{X}} \circ \hat{\mathbf{Y}}) \Phi_c^\top\end{aligned}$$

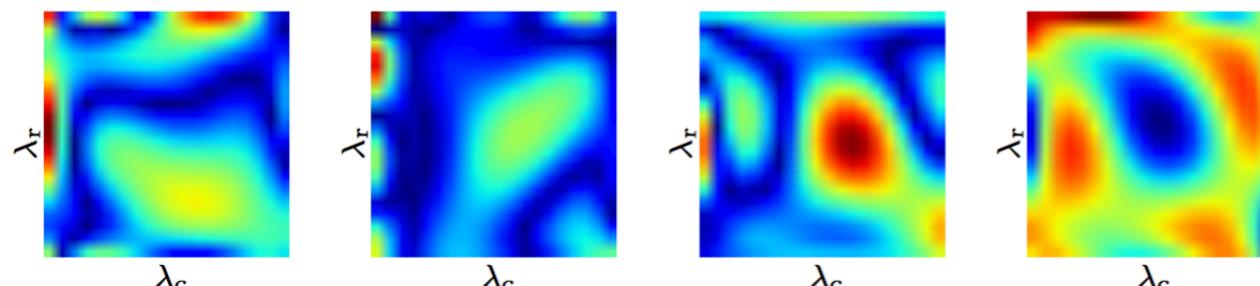
[25] Monti, Bresson, Bronstein 2017

# Multi-graph spectral filters

- Parametrize the multi-graph filter as a smooth function of column- and row frequencies w/ e.g. bivariate Chebyshev polynomial

$$w_{\Theta}(\tilde{\lambda}_c, \tilde{\lambda}_r) = \sum_{j,j'=1}^r \theta_{jj'} T_j(\tilde{\lambda}_c) T_{j'}(\tilde{\lambda}_r)$$

where  $\Theta = (\theta_{jj'})$  is matrix of coefficients and  $-1 \leq \tilde{\lambda}_c, \tilde{\lambda}_r \leq 1$



Examples of multi-graph spectral filters (shown is  $|\tau(\tilde{\lambda}_c, \tilde{\lambda}_r)|$ )

- Multi-graph convolutional layer: Apply filter to matrix X

$$\mathbf{Y} = \sum_{j,j'=1}^r \theta_{jj'} T_j(\tilde{\Delta}_c) \mathbf{X} T_{j'}(\tilde{\Delta}_r)$$

where the scaled Laplacians  $\tilde{\Delta}_c = 2\lambda_{c,n}^{-1}\Delta_c - \mathbf{I}$  and  $\tilde{\Delta}_r = 2\lambda_{r,m}^{-1}\Delta_r - \mathbf{I}$

# Multi-graph ConvNets<sup>[25]</sup>

- Multi-graph spectral filters

$$\mathbf{Y}_l^{(k)} = \xi \left( \sum_{l'=1}^p \sum_{j,j'=0}^r \theta_{jj' ll'} T_j(\tilde{\Delta}_r) \mathbf{Y}_{l'}^{(k-1)} T_{j'}(\tilde{\Delta}_c) \right) \quad \begin{array}{l} l = 1, \dots, q \\ l' = 1, \dots, p \end{array}$$

applied to  $p$  input channels ( $m \times n$  matrices  $\mathbf{Y}_1^{(k-1)}, \dots, \mathbf{Y}_p^{(k-1)}$ ) and producing  $q$  output channels  $\mathbf{Y}_1^{(k)}, \dots, \mathbf{Y}_q^{(k)}$

- ☺ Filters have guaranteed  $r$ -hops support on both graphs
- ☺  $O(1)$  parameters per layer
- ☹  $O(nm)$  computational complexity (assuming sparse graphs)

[25] Monti, Bresson, Bronstein 2017

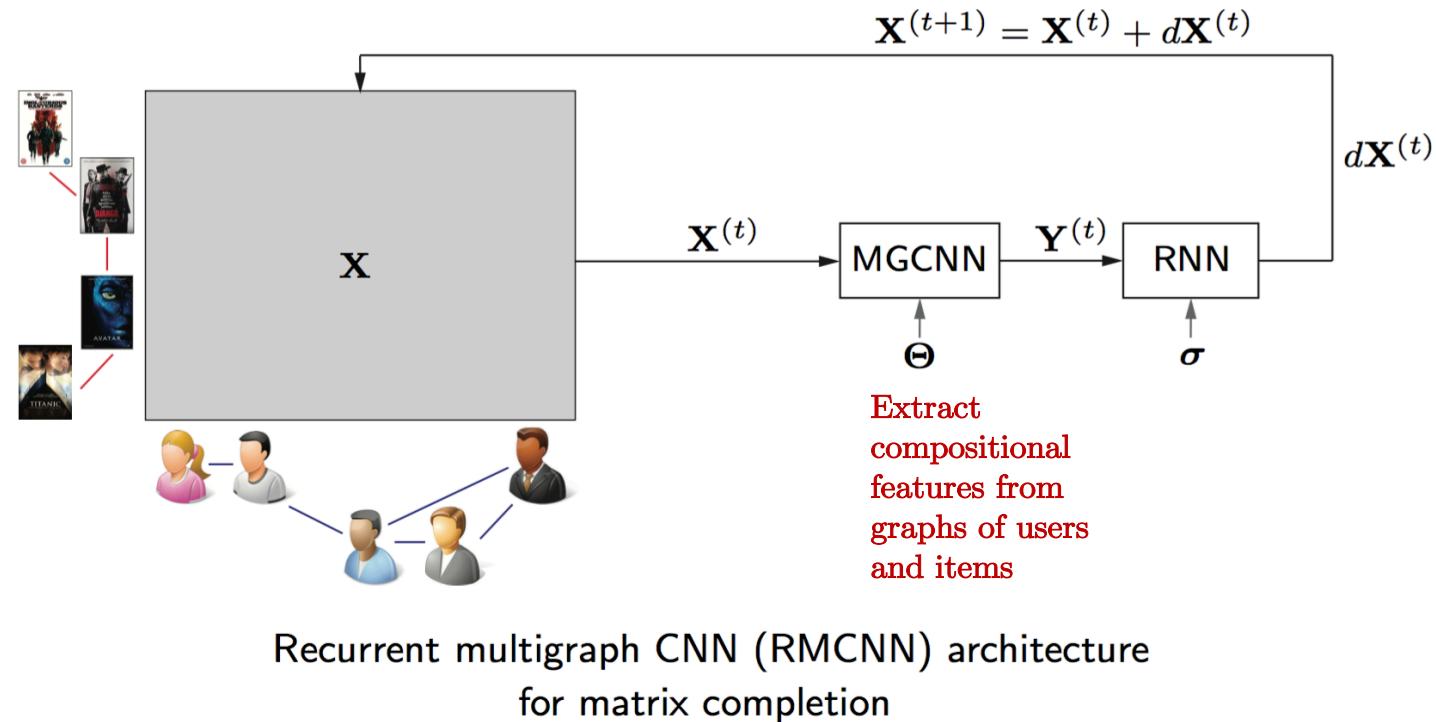
# Factorized multi-graph ConvNets

- Two spectral convolutional layers applied to each of the factors  $\mathbf{W}$ ,  $\mathbf{H}$

$$\begin{aligned}\mathbf{u}_l &= \xi \left( \sum_{l'=1}^p \sum_{j=0}^r \theta_{\text{r},jl} l' T_j(\tilde{\Delta}_{\text{r}}) \mathbf{w}_{l'} \right) & l &= 1, \dots, q \\ \mathbf{v}_l &= \xi \left( \sum_{l'=1}^p \sum_{j'=0}^r \theta_{\text{c},j'l} l' T_{j'}(\tilde{\Delta}_{\text{c}}) \mathbf{h}_{l'} \right) & l' &= 1, \dots, p\end{aligned}$$

- ☺ Filters have guaranteed  $r$ -hops support on both graphs
- ☺  $O(1)$  parameters per layer
- ☺  $O(n+m)$  computational complexity (assuming sparse graphs)
- ☹ May loose some coupling information

# Matrix completion with Recurrent Multi-Graph ConvNets<sup>[25]</sup>

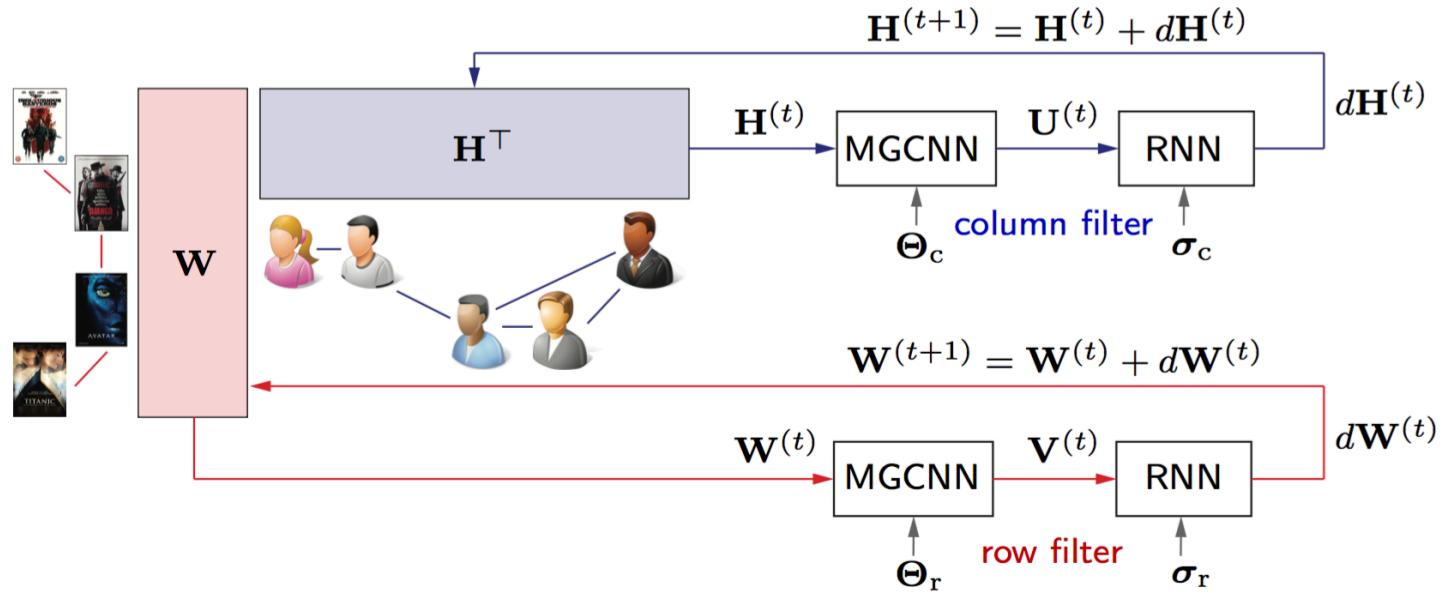


$$\min_{\Theta, \sigma} \|\mathbf{X}_{\Theta, \sigma}^{(T)}\|_{\mathcal{G}_r}^2 + \|\mathbf{X}_{\Theta, \sigma}^{(T)}\|_{\mathcal{G}_c}^2 + \frac{\mu}{2} \|\Omega \circ (\mathbf{X}_{\Theta, \sigma}^{(T)} - \mathbf{Y})\|_F^2$$

Complexity:  $O(nm)$

[25] Monti, Bresson, Bronstein 2017

# Factorized version



Separable recurrent multigraph CNN (sRMCNN) architecture  
for matrix completion in factorized form

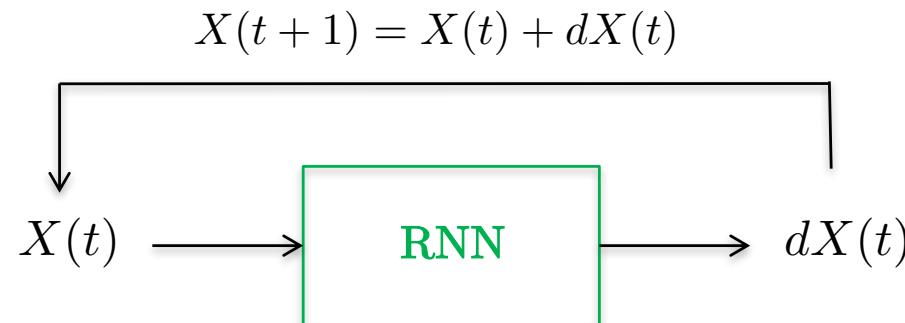
$$\min_{\theta_r, \sigma_r, \Theta_c, \sigma_c} \|\mathbf{W}_{\theta_r, \sigma_r}^{(T)}\|_{\mathcal{G}_r}^2 + \|\mathbf{H}_{\theta_c, \sigma_c}^{(T)}\|_{\mathcal{G}_c}^2 + \frac{\mu}{2} \|\Omega \circ (\mathbf{W}_{\theta_r, \sigma_r}^{(T)} (\mathbf{H}_{\theta_c, \sigma_c}^{(T)})^\top - \mathbf{Y})\|_F^2$$

Complexity:  $O(n+m)$

[25] Monti, Bresson, Bronstein 2017

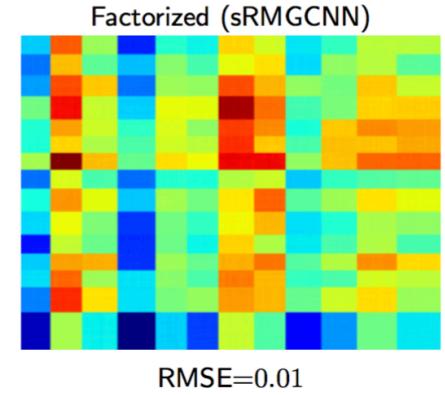
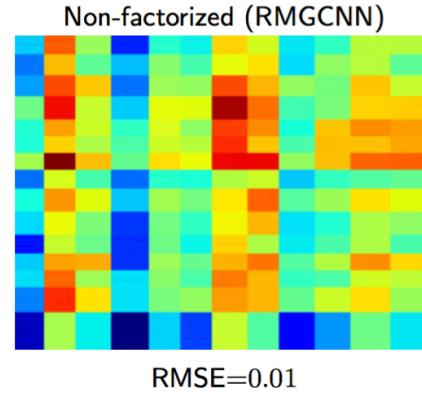
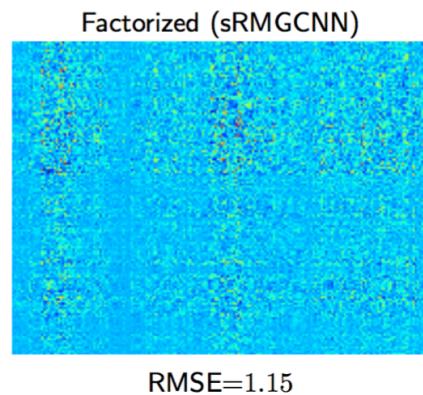
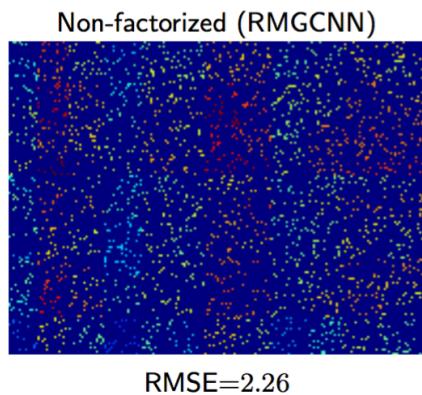
# Incremental updates with RNNs

- The RNN/LSTM<sup>[26]</sup> cast the completion problem as **non-linear diffusion** process.
  - Replace multi-layer CNNs by simple RNN + diffusion  
⇒ **much less parameters to learn** (reduce overfitting).  
⇒ **more favorable for highly sparse entries** (better info diffusion).
  - Fix #parameters of the model **independently of entry sparsity**.
  - Allows to **learn temporal dynamics of entries** if such information is available.



[26] Hochreiter, Schmidhuber 1997

# Matrix completion methods comparison: synthetic



Matrix completion results on a synthetic dataset (initialization)

Matrix completion results on a synthetic dataset after  $t = 10$  instances

| Method                     | #Params  | Complexity                | RMSE          |
|----------------------------|----------|---------------------------|---------------|
| GMC <sup>1</sup>           | $mn$     | $mn$                      | 0.3693        |
| GRALS <sup>2</sup>         | $m + n$  | $m + n$                   | 0.0114        |
| <b>sRMGCNN<sup>3</sup></b> | <b>1</b> | <b><math>m + n</math></b> | <b>0.0106</b> |
| <b>RMGCNN<sup>3</sup></b>  | <b>1</b> | <b><math>mn</math></b>    | <b>0.0053</b> |

Comparison of geometric matrix completion methods on synthetic data  
using both user+movie graphs

Methods: <sup>1</sup>Kalofolias et al. 2014; <sup>2</sup>Rao et al. 2015; <sup>3</sup>Monti, Bresson, Bronstein 2017

# Matrix completion methods comparison: real

| Method                              | MovieLens <sup>1</sup> | Flixster <sup>2</sup> | Douban <sup>3</sup> | Yahoo <sup>4</sup> |
|-------------------------------------|------------------------|-----------------------|---------------------|--------------------|
| IMC <sup>5</sup>                    | 1.653                  | –                     | –                   | –                  |
| GMC <sup>6</sup>                    | 0.996                  | –                     | –                   | –                  |
| MC <sup>7</sup>                     | 0.973                  | –                     | –                   | –                  |
| GRALS <sup>8</sup>                  | 0.945                  | 1.245                 | 0.833               | 38.042             |
| <b>sRGCNN (Cheb)<sup>9</sup></b>    | <b>0.929</b>           | <b>0.926</b>          | <b>0.801</b>        | <b>22.415</b>      |
| <b>sRGCNN (Cayley)<sup>10</sup></b> | <b>0.922</b>           | –                     | –                   | –                  |

Accuracy (RMS error) of matrix completion methods on real data

Data: <sup>1</sup>Miller et al. 2003; <sup>2</sup>Jamali, Ester 2010; <sup>3</sup>Ma et al. 2011; <sup>4</sup>Dror et al. 2012

Methods: <sup>5</sup>Jain, Dhillon 2013; <sup>6</sup>Kalofolias et al. 2014; <sup>7</sup>Candès, Recht 2012; <sup>8</sup>Rao et al. 2015; <sup>9</sup>Monti, Bresson, Bronstein 2017; <sup>10</sup>Levie et al. 2017

# Outline

- Deep Learning
  - *A brief history*
  - *High-dimensional learning*
- Convolutional Neural Networks
  - *Architecture*
  - *Non-Euclidean Data*
- Spectral Graph Theory
  - *Graphs and operators*
  - *Spectral decomposition*
  - *Fourier analysis*
  - *Convolution*
  - *Coarsening*
- Spectral ConvNets
  - *SplineNets*
  - *ChebNets\* [NIPS'16]*
  - *GraphConvNets*
  - *CayleyNets\**
- Spectral ConvNets on Multiple Graphs
  - *Recommender Systems\* [NIPS'17]*
- Conclusion

# Related works

- Classes of graph ConvNets:
  - (1) Spectral approaches
  - (2) Spatial approaches
- Spatial approaches
  - Geodesic CNN: Masci, Boscaini, Bronstein, Vanderghenst 2015
  - Anisotropic CNN: Boscaini, Masci, Rodola, Bronstein 2016
  - MoNet: Monti, Boscaini, Masci, Rodola, Svoboda, Bronstein 2017
- Graph Neural Networks:
  - Gori, Monfardini, Scarselli 2005
  - Li, Tarlow, Brockschmidt, Zemel 2015
  - Sukhbaatar, Szlam, Fergus 2016
  - Bruna, Li 2017
- Several other works

# Conclusion

- Contributions:
  - Generalization of ConvNets to data on graphs (fixed)
  - Exact/tight localized filters on graphs
  - Linear complexity for sparse graphs
  - GPU implementation
  - Rich expressivity
  - Multiple graphs
- Several potential applications in network sciences.

# Papers & code

- **Papers**

- Convolutional neural networks on graphs with fast localized spectral filtering, M Defferrard, X Bresson, P Vandergheynst, NIPS, 2016, arXiv:1606.09375
- Geometric matrix completion with recurrent multi-graph neural networks, F Monti, MM Bronstein, X Bresson, NIPS, 2017, arXiv:1704.06803
- CayleyNets: Graph Convolutional Neural Networks with Complex Rational Spectral Filters, R Levie, F Monti, X Bresson, MM Bronstein, 2017, arXiv:1705.07664

- **Code**

- [https://github.com/xbresson/graph\\_convnets\\_pytorch](https://github.com/xbresson/graph_convnets_pytorch)

## Graph ConvNets in PyTorch

September 30, 2017

### Xavier Bresson

 <http://www.ntu.edu.sg/home/xbresson>  
 <https://github.com/xbresson>  
 <https://twitter.com/xbresson>



### Description

Prototype implementation in PyTorch of the NIPS'16 paper:  
Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering  
M Defferrard, X Bresson, P Vandergheynst  
Advances in Neural Information Processing Systems, 3844-3852, 2016  
ArXiv preprint: [arXiv:1606.09375](https://arxiv.org/abs/1606.09375)

### Code objective

The code provides a simple example of graph ConvNets for the MNIST classification task.  
The graph is a 8-nearest neighbor graph of a 2D grid.  
The signals on graph are the MNIST images vectorized as \$28^2 \times 1\$ vectors.

### Installation

```
git clone https://github.com/xbresson/graph_convnets_pytorch.git
cd graph_convnets_pytorch
pip install -r requirements.txt # installation for python 3.6.2
python check_install.py
jupyter notebook # run the 2 notebooks
```

# Tutorials



## Computer Vision and Pattern Recognition (CVPR)

Geometric deep learning on graphs and manifolds

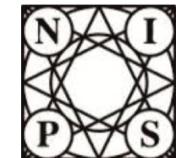
M. Bronstein, J. Bruna, A. Szlam, X. Bresson, Y. LeCun

July 21, 2017

Slides: [https://www.dropbox.com/s/appafg1fumb6u7f/tutorial\\_CVPR17\\_DL\\_graphs.pdf?dl=0](https://www.dropbox.com/s/appafg1fumb6u7f/tutorial_CVPR17_DL_graphs.pdf?dl=0)

Video will be posted on YouTube/ComputerVisionFoundation Videos

## Neural Information Processing Systems (NIPS)



Geometric deep learning on graphs and manifolds

M. Bronstein, J. Bruna, A. Szlam, X. Bresson, Y. LeCun

Dec 4, 2017

# Workshop

## New Deep Learning Techniques

Feb 5-9, 2018

### Organizers

Xavier Bresson, NTU

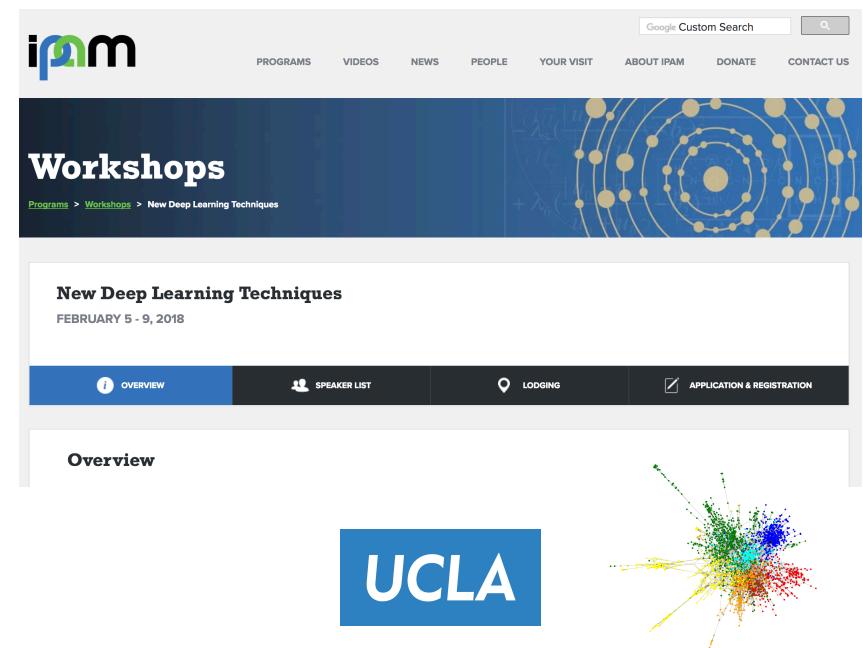
Michael Bronstein, USI/Intel

Joan Bruna, NYU/Berkeley

Yann LeCun, NYU/Facebook

Stanley Osher, UCLA

Arthur Szlam, Facebook



<http://www.ipam.ucla.edu/programs/workshops/new-deep-learning-techniques>



Thank you

 <http://www.ntu.edu.sg/home/xbresson>

 <https://github.com/xbresson>

 <https://twitter.com/xbresson>